

Heterogeneous Expert PLE: An Explainable Multi-Task Architecture for Financial Product Recommendation

Seonkyu Jeong¹, Euncheol Sim¹, Youngchan Kim¹

¹*Independent Research*

Corresponding author: Seonkyu Jeong (ORCID: 0009-0005-3291-9112)

Abstract. Financial product recommendation requires not only predictive accuracy but also regulatory-compliant explainability. Existing multi-task learning (MTL) approaches (MMoE, PLE) are validated on 2–4 homogeneous tasks, leaving open whether they scale to production-scale heterogeneous workloads where binary, multiclass, and regression tasks compete for shared capacity. We present three contributions. **First**, we propose *Heterogeneous Expert PLE*, where seven architecturally distinct experts (DeepFM, Temporal Ensemble, Hyperbolic GCN, PersLay, Causal, LightGCN, Optimal Transport) share a common basket with a FeatureRouter assigning each expert its designated feature groups, providing a structural guarantee against expert collapse and inherent explainability through business-interpretable gate weights. **Second**, ablation on a 12-task benchmark (7 binary + 2 multiclass + 3 regression, 1M customers) shows that our *loss-level* inter-task transfer mechanism (a TAG + GradNorm hybrid retained under the “adaTT” label, distinct from the representation-level adaTT of [22]) has a **null effect** at this scale: after correcting five implementation bugs inherited from the on-prem port, the AUC difference with vs. without adaTT shrinks to -0.001 , indistinguishable from single-seed measurement noise. The single largest improvement instead comes from correcting a subtle uncertainty-weighting implementation gap where per-task loss weights were silently ignored. **Third**, we experiment with *GradSurgery*, a task-type gradient projection method that replaces loss-level transfer with gradient-level conflict resolution, reducing the 132 task-pair problem to 3 task-type-group projections. However, GradSurgery shows no meaningful advantage over the PLE-only baseline while requiring significantly more VRAM (due to the retained computation graph), and is not adopted for production. In the ablation study, PLE with softmax gating achieves the best NDCG@3 (0.714) among gate variants. Softmax also outperforms sigmoid in this heterogeneous setting, reversing findings from the homogeneous-task literature. Both GradSurgery (AUC 0.673, F1-macro 0.203) and the corrected adaTT implementation (sigmoid+adaTT AUC 0.672) maintain baseline performance; neither provides a gain. The operational motivation — consolidating 12 individual models into a single MTL model for unified training, serving, and monitoring — is validated: the shared-bottom baseline already exceeds per-task XGBoost ceilings (on the synthetic benchmark; in a clean OOT measurement on operational real data, a raw-feature LGBM outperforms the teacher on 10 of 11 active tasks, so the operational value of consolidation lies in unified training, representation, and explanation rather than per-task predictive superiority — see Section 6.4), and PLE provides additional gains on ranking metrics without sacrificing classification performance.

Keywords: Multi-task learning, Progressive Layered Extraction, Mixture of Experts, Financial recommendation, Explainable AI, Ablation study

1 Introduction

1.1 Problem Definition

Financial product recommendation differs fundamentally from e-commerce or content recommendation. The primary deliverable is not a probability score but a *reason that the customer can accept*. Three audiences must be persuaded: customers (“Why this product for me?” — trust leads to conversion), relationship managers (“Why recommend this to this customer?” — sales justification), and regulators (Korea’s FSC Financial-Sector AI Guideline [11], EU AI Act [8]) (“Why was this decision made?” — compliance obligation). Korea’s AI Basic Act [29]

further classifies financial recommendation as potentially high-impact AI.

Existing approaches fall short on this persuasion requirement. Single-task models cannot jointly predict churn, product affinity, and customer lifetime value [3]. Model ensembles (N separate models) multiply management overhead and serving cost, and “MLP #3 contributed 28%” provides no business meaning. Post-hoc explanations (SHAP, LIME) are decoupled from model internals, computationally expensive at serving time, and demonstrably unstable under input perturbation [24, 36, 37].

1.2 Core Insight

Conventional recommendation models function as black-box score generators — features go in, probabilities come out, with only statistical correlations as justification. As Pearl [34] argues, there is a fundamental gap between *seeing* (association) and *understanding* (causation). Humans are not persuaded by correlations; they require *causal narratives* [33].

If each expert in a multi-task model captures a *different kind of “why”* — temporal trends, hierarchical product structure, causal pathways, collaborative patterns — then the gating mechanism itself becomes an explanation: “This recommendation is driven primarily by your spending trend (Temporal, 35%) and product category fit (HGCN, 28%).”¹

This principle motivates the architecture proposed in this work.

The key realization is that financial recommendation is not about a single prediction (“will this customer buy?”) but about **understanding a customer as a whole person** from multiple perspectives simultaneously: Will they churn? (lifecycle). How much will they spend? (value). What product fits them? (consumption). How do they behave? (engagement). Each question demands a different analytical lens — temporal analysis cannot answer hierarchical product fit, and graph-based collaborative signals cannot predict spending trajectories. This is why heterogeneous experts are not merely a performance optimization but a structural necessity for multifaceted customer understanding.

This perspective aligns with a broader shift in both academia and regulation. Pearl’s *Ladder of Causation* [33] distinguishes three levels: association, intervention, and counterfactuals. Most recommendation systems operate at level 1 (association); our architecture operates primarily at level 1 (association), with the Causal expert approaching level 2 (intervention) for tasks where its gate weight is dominant. Regulatory frameworks — the EU AI Act, Korea’s FSC Financial-Sector AI Guideline (in force June 2026) and AI Basic Act — demand this shift toward structurally transparent explanations (detailed regulatory analysis in Section 2.4).

1.3 Contributions

1. **Heterogeneous Shared Expert Basket with Structural Collapse Guarantee:** We replace PLE’s homogeneous MLP experts with seven architecturally distinct experts (DeepFM, Temporal

Ensemble, HGCN, PersLay, Causal, LightGCN, Optimal Transport). Unlike prior “heterogeneous” MoE work that varies expert *size* [43] or *modality* [1], we vary the fundamental *inductive bias*, providing a structural guarantee against expert collapse — a persistent failure mode in homogeneous MoE/PLE deployments [21].

2. **FeatureRouter: Heterogeneous Architecture × Heterogeneous Input:** Beyond architectural diversity, each expert receives only its designated feature groups (declared via expert routing declarations at the feature group level), not the full 1211D input. This “heterogeneous architecture × heterogeneous input” design eliminates irrelevant features per expert (per-expert dims: 32D–977D), reducing model parameters substantially (V1 baseline at 349D schema: 2.70M measured; V2 1211D schema: 0.97M across the seven routed shared experts vs. 2.96M for the same basket with routing disabled, measured) while strengthening each expert’s specialization.
3. **Inherent Explainability:** Because each expert encodes a named mathematical operation (not a generic MLP), Customized Gate Control (CGC) gate weights directly yield business-interpretable explanations without post-hoc attribution methods. Crisp attribution to 1–2 experts, however, is limited to tasks with a dominant expert; most tasks default to multi-factor explanation (Section 5.5, Section 6.4).
4. **Multi-disciplinary Feature Engineering:** Features derived from eleven academic disciplines — including unconventional applications of chemical kinetics (spending activation rate), epidemic modeling (product adoption diffusion), criminological Routine Activity Theory (transaction regularity), and wave interference (spending periodicity) — serve dual roles as learning signals and recommendation context that is reverse-mapped to business language for customer-facing explanations.
5. **Financial DNA Task Grouping:** Four task groups (engagement, lifecycle, value, consumption) aligned with the two-axis decomposition (Financial DNA × Data Modality). Expert selection itself operates *per-task* via the CGC gate (12 independent routing decisions); the task groups provide semantic structure for post-hoc interpretation and for the rule-based fallback layer’s template selection rather than gating the router directly.

¹This three-expert decomposition is an idealized example for tasks with a dominant expert. On operational data the gate is more uniform (normalized entropy 0.67–0.97), so such crisp decompositions materialize only for the few tasks with a dominant expert (Section 6.4).

6. **Gate Type Analysis for Heterogeneous MTL:** We demonstrate that softmax gating outperforms sigmoid in 12-task heterogeneous settings (7 binary + 2 multiclass + 3 regression), reversing the conventional preference from homogeneous-task literature [30, 41]. The reversal is attributed to softmax’s protective isolation of minority-type tasks from majority-type gradient corruption.
7. **Loss-Level vs. Gradient-Level Transfer:** Our loss-level inter-task transfer mechanism (a TAG [10] + GradNorm [6] hybrid, retained under the internal “adaTT” label but distinct from the representation-level adaTT of [22] — see Section 3.6) has a null effect at 12-task scale after correcting five implementation bugs: the AUC change with adaTT vs. without is -0.001 , within single-seed measurement noise. An earlier draft reported -0.019 , attributed at the time to algorithmic instability; the corrected measurements locate the cause in the implementation rather than in the loss-level mechanism per se (see Section 5.4 table caption). We additionally experiment with GradSurgery (gradient-level projection between 3 task-type groups), which also shows no gain and incurs VRAM overhead from the retained computation graph; GradSurgery is not adopted for production either.
8. **Uncertainty Weighting Correction:** We identify and fix a subtle implementation gap where per-task loss weights were silently ignored under uncertainty weighting — yielding the single largest performance improvement ($+0.018$ NDCG@3, $+0.031$ F1-macro), larger than any architectural change.
9. **Comprehensive Ablation:** 23 scenarios (14 joint feature+expert + 9 structure cross) on a reproducible 1M-customer benchmark with Gaussian Copula + latent variable variance budget.
10. **Config-driven Pipeline:** End-to-end system (feature engineering $\rightarrow$ training $\rightarrow$ distillation $\rightarrow$ serving) controlled by a split-config of three YAML files (`pipeline.yaml` for common model/training/distillation/AWS settings, `feature_groups.yaml` for feature-group routing, and `configs/datasets/{name}.yaml` for dataset-specific tasks/labels/ablation), enabling deployment by teams with 1–2 ML engineers.
11. **Reproducible Benchmark:** Synthetic data generation with variance budget [32] for controlled AUC ceilings, validated against XGBoost baselines, with open-source code and fixed seeds.

2 Related Work

2.1 Multi-Task Learning for Recommendation

The progression from Shared-Bottom [3] to MMoE [25] to PLE [41] reflects increasing sophistication in managing negative transfer between tasks. MMoE introduced per-task gating over a shared expert pool, while PLE further separated shared and task-specific experts with a progressive extraction structure. AdaTT [22] extended this with adaptive inter-task transfer at the **representation level** (per-task fusion over expert outputs via learned gating). Gradient-based task-affinity methods such as Task Affinity Groupings [10] and adaptive loss balancing [6] are a parallel line of work that modulates **loss scalars** using per-task gradient statistics rather than mixing representations; our implementation (Section 3.6) follows this second line and should not be confused with the representation-level mechanism of [22], despite our retention of the “adaTT” label for continuity.

Other notable MTL architectures include ESMM [26] (entire-space modeling for conversion rate), STAR [39] (star topology for multi-domain CTR), and M3oE [44] (multi-domain multi-task MoE with AutoML structure search).

However, all prior MoE/PLE architectures employ **homogeneous experts** — multiple MLPs with identical architecture but different initializations. This limits the diversity of learned representations to what parameter variation alone can achieve, and leads to a well-documented failure mode: **expert collapse**.

2.2 Mixture of Experts and the Expert Collapse Problem

The MoE paradigm [38] and its successors (Switch Transformer [9]) demonstrate the power of conditional computation. However, homogeneous MoE architectures suffer from **expert collapse**, where all gates converge to routing inputs to a single expert, effectively reducing the MoE to a shared-bottom architecture. Pinterest Engineering reported that their MMoE deployment “would oftentimes collapse into a state where all tasks use the same single expert.” Kuaishou’s HoME [21] confirmed that “expert collapse occurs when all gates assigned larger weights to a single shared expert and almost ignored other shared experts.”

Existing mitigations operate post-hoc: Gram-Schmidt orthogonalization to force diverse expert representations, expert normalization to align output distributions, and load-balancing losses to distribute routing. These methods mitigate the symptom but do not eliminate the

structural source of the problem — homogeneous experts have no inherent reason to specialize differently because their only source of diversity is random initialization.

Recent work has explored limited forms of heterogeneity: Jamba [1] combines Transformer and Mamba layers within MoE for LLMs (2 architecture types), and MOWST [43] pairs lightweight MLP “weak experts” with GNN “strong experts” in a hierarchical arrangement (2 types, unequal roles). M3oE [44] separates experts by domain but uses identical MLP architectures within each domain. MoE++ introduces computation-level heterogeneity (zero-compute vs. FFN experts) but not architectural heterogeneity.

Furthermore, recent theoretical work demonstrates that the standard softmax gating mechanism itself contributes to representation collapse. Nguyen et al. [30] prove that sigmoid gating achieves higher sample efficiency than softmax gating in MoE models: softmax’s sum-to-one constraint forces experts into unnecessary competition, where increasing one expert’s weight requires decreasing another’s, leading to representation collapse even when experts provide complementary signals. Sigmoid gating eliminates this inter-expert competition by allowing each expert’s contribution to be evaluated independently, achieving faster convergence and better utilization of all experts. This finding is especially relevant when experts are *heterogeneous*: since their outputs have fundamentally different scales and distributions (e.g., hyperbolic embeddings vs. persistence diagrams), softmax competition is particularly harmful — suppressing a topological expert to increase a temporal expert’s weight means losing irreplaceable topological signal, not merely a redundant representation.

Gap: No prior work composes experts with **fundamentally different inductive biases** — graph convolution, state-space models, topological persistence, causal DAG constraints, optimal transport — as equal peers within a shared expert basket. Expert collapse manifests in two forms: (1) *function space collapse* — homogeneous experts converge to identical functions due to shared architecture and gradient dynamics; (2) *routing collapse* — gate weights concentrate on a single expert, effectively deactivating others regardless of their functional diversity. Our heterogeneous expert design structurally prevents (1). For (2), prior work suggests sigmoid gating [30] mitigates routing collapse through independent expert contributions. However, our ablation (Section 5.4) reveals that in heterogeneous task settings, softmax’s competitive isolation better protects minority-type tasks — a reversal of the homogeneous-task finding.

Our work provides a **structural guarantee** against expert collapse: a DeepFM expert is far less likely to converge to the same functional behavior as an HGCN expert, because the two architectures encode fundamentally different inductive biases and mathematical operations.

2.3 Explainability in Recommendation

SHAP [24] and LIME [36] provide model-agnostic explanations but suffer from instability, computational cost, and disconnect from model internals. Integrated Gradients [40] offers theoretically grounded attribution but still operates post-hoc.

Gap: Structural (inherent) explainability that produces business-meaningful explanations as a natural byproduct of the forward pass, without additional computation.

2.4 Financial Recommendation and Regulation

Recent work on deep learning for financial product recommendation [5, 28] demonstrates improvements over collaborative filtering, while Martinez-Plumed et al. [27] show that sequential models on longitudinal transaction data improve cross-selling prediction. However, none of these systems address regulatory explainability requirements.

The EU AI Act [8] classifies financial recommendation as high-risk AI, mandating transparency (Art. 13), human oversight (Art. 14), and robustness (Art. 15). Korea’s FSC Financial-Sector AI Guideline [11] (in force June 2026) and AI Basic Act [29] impose similar requirements. In particular, the guideline’s Reliability principle (§4.4) distinguishes global and local explanation and, where a legal duty to explain applies, recommends XAI at the SHAP level or above — which our inherent explainability addresses by producing both gate-weight-based global explanations and per-prediction local explanations. Current approaches rely on post-hoc SHAP/LIME for regulatory compliance, which has documented limitations in financial contexts [37].

Gap: No recommendation system provides verifiable mapping from regulatory requirements to architecture components, with explanations generated structurally rather than post-hoc.

DNA	Question	Tasks
Engagement	What do they <i>do</i> ?	next_mcc mcc_diversity_trend top_mcc_shift
Lifecycle	Where <i>are</i> they?	churn_signal segment_prediction
Value	How much are they <i>worth</i> ?	cross_sell_count product_stability
Consumption	What <i>will</i> they buy?	will_acquire_* (5) nba_primary ²

Table 1: Financial DNA decomposition (Axis 1). Four tasks removed in early development (income tier, tenure stage, spend level, engagement score) were deterministic feature transformations.

3 Architecture

3.1 Reductionist Framework: Two Axes of Decomposition

Before describing the architecture, we present the analytical framework that governs all subsequent design decisions.

The complexity of financial customer understanding is decomposed along two orthogonal axes:

Axis 1: Who is the customer? (Financial DNA)

We decompose the question “who is this customer?” into four irreducible dimensions, each capturing a fundamentally different aspect of customer identity:

Axis 2: What form does the information take? (Data Modality)

Customer data exists in structurally distinct modalities, each requiring a different mathematical tool to extract meaningful signals:

The cross-product of these two axes defines the architecture: each (DNA $\times$ Modality) cell determines which expert processes which features for which task group. The 7 heterogeneous experts, 17 feature groups, and 4 task groups are not arbitrary design choices but necessary consequences of this two-axis decomposition. A homogeneous MLP expert basket ignores Axis 2 entirely — it treats hierarchical, temporal, and topological data identically — which is why it cannot achieve the task-type-specific specialization that our ablation demonstrates.

3.2 Design Philosophy

The architecture emerged from severe real-world constraints at a Korean public financial institution. A team of three — one data scientist serving as PM plus two

Modality	Examples	Expert	Generator
State	is_active, holdings	DeepFM	base
Snapshot	balance, demographics	GMM, DeepFM	GMM clustering
Short-term series	recent transactions	Transformer	temporal
Long-term series	monthly trends	Mamba	mamba temporal
Disrupted series	dormant→active	LNN	model derived
Hierarchy	MCC merchant category tree	HGCN	merchant_hierarchy
Relations	customer-product graph	LightGCN	graph collab.
Topology	behavioral shape	PersLay	TDA global/local
Causality	behavioral causation	Causal	causal features

Table 2: Data modality decomposition (Axis 2). Each modality is mapped to an expert and feature generator.

²nba_primary targets *product groups* (7 classes: no NBA, savings/guarantee, checking, deposits, investments, credit loans, debits), not individual product indices. This redesign from 25 individual-product classes achieves better class balance and enables ranking-based evaluation via NDCG@K.

engineers — needed to replace a legacy ALS-based collaborative filtering system with a next-generation recommendation model. The PM holds the GARP Financial Risk Manager (FRM) certification, with a career spanning credit and market risk analysis, regulatory compliance, MyData licensing, big data platform operations, data science projects, and recommendation system management — a background that directly informed the architectural and compliance decisions described in this paper.

The constraints were formidable: no dedicated ML infrastructure budget, a single desktop-grade GPU (NVIDIA RTX 4070, 12GB VRAM) as the only training hardware, no GPU inference servers for deployment, and strict regulatory requirements (Korea’s FSC Financial-Sector AI Guideline, EU AI Act).

Rather than treating these constraints as limitations, the team adapted its methodology at every level:

- **AI-augmented development** using Claude (Anthropic), Gemini (Google), and Cursor, with each team member leading a parallel team of AI agents;³
- **Parameter-efficient architecture design** where structural inductive biases replace the brute-force capacity of large MLPs;
- **Knowledge distillation** to LGBM for GPU-free CPU inference on AWS Lambda;
- **Config-driven pipeline** requiring only a split-config of three YAML files to control the entire system, enabling operation by a minimal team.

Early exploration considered a Black-Litterman-inspired approach, treating multiple models’ predictions as “expert views” combined via Bayesian updating. This was abandoned because the Bayesian combination process obscured each model’s contribution, making business-meaningful explanation impossible — a critical failure for regulatory compliance.

The key reframing was: instead of combining *models* externally, combine *experts* internally within a single model. This led to the selection of PLE [41] with a critical modification: replacing homogeneous MLP experts with structurally heterogeneous experts, each encoding a different inductive bias.

3.2.1 Structural Isomorphism: Why Cross-Disciplinary Methods Work

The multi-disciplinary feature engineering in this work is not arbitrary borrowing from unrelated fields. It is grounded in *structural isomorphism* — the observation

that mathematically equivalent structures recur across different domains, enabling tools developed in one field to be rigorously applied to structurally equivalent problems in another.

Consider: a dormant customer resuming spending follows the same dynamics as a chemical reaction overcoming an activation energy barrier. Product adoption spreading through a customer network follows the same compartmental dynamics as an SIR epidemic model. These are not metaphors — the governing equations are identical, and the solutions transfer with mathematical rigor. We emphasize that these are not loose analogies: the governing differential equations are structurally identical (see Appendix E for side-by-side comparison). The validity of each mapping can be independently verified by checking whether the mathematical preconditions of the source domain’s theorems hold in the financial domain.

This principle has a distinguished history in science and engineering: Shannon’s information entropy was directly borrowed from Boltzmann’s thermodynamic entropy; the Black-Scholes option pricing model is derived from the heat diffusion equation; Google’s PageRank is an application of Markov chain stationary distributions. In each case, recognizing structural isomorphism between two domains enabled a breakthrough that neither domain could have achieved alone.

Our contribution is applying this principle *systematically* to financial customer understanding: for each aspect of customer behavior, we identify a domain whose mathematical tools are structurally isomorphic to the underlying dynamics, and instantiate that domain’s methods as both features (for prediction) and expert architectures (for interpretation). The multi-disciplinary approach is not a collection of ad-hoc features but a principled application of analogical reasoning through structural equivalence.

Four principles guide the architecture:

1. **Robust Explainability:** Model structure itself generates explanations. Each expert’s name carries business meaning (e.g., “Temporal” = time patterns, “HGCN” = product hierarchy), so CGC gate weights are directly interpretable without post-hoc methods.
2. **Expert Contribution Diagnostics:** The heterogeneous design enables per-expert contribution analysis — identifying which inductive biases are beneficial vs. harmful in a given data regime. Homogeneous experts would show uniform, uninformative degradation.

³The guideline’s Auxiliary-means principle (§3) cites Anthropic’s Claude Mythos system card as a frontier-AI governance example [11]. This note adds context for the tool-use fact and is not a compliance claim.




Figure 1: Heterogeneous Expert PLE architecture overview.

3. **Flexible Extensibility:** New features, tasks, or experts are added via YAML configuration, not code changes. The system has been extended from 4 to 12 tasks without architectural modification.⁴
4. **Unified Manageability:** The entire pipeline (feature engineering → training → distillation → serving → monitoring) is controlled by a split-config of three YAML files (`pipeline.yaml` + `feature_groups.yaml` + `configs/datasets/{name}.yaml`), reducing operational overhead for teams with limited ML engineering resources.

3.3 Data Axis Classification

Financial customer data exhibits inherently multi-modal structure. We classify data along multiple axes, each mapped to an optimal feature generator and expert:

The complete data-axis-to-expert-to-feature-generator mapping is shown in Table 2.

Note: Short-term, Long-term, and Disrupted series map to the three sub-components of the Temporal Ensemble expert (Transformer, Mamba, LNN respectively). GMM is a feature generator, not a standalone expert.

3.4 Heterogeneous Expert Basket

Unlike standard PLE where shared experts are identical MLPs, our shared expert basket contains seven structurally distinct networks (Table 3). Each expert corre-

sponds to one or more modality axes (Table 2), ensuring that every structurally distinct data type is processed by an architecture designed for it:

Because each expert has a named inductive bias, the gate weights tell the customer *why* in business terms — not “hidden unit 47 activated” but “your spending trend drove this recommendation.” Each expert was selected based on a specific gap in financial customer understanding that no other expert type addresses:

- **DeepFM** [14]: Financial behavior is driven by feature interactions (e.g., income × product holdings × channel preference). FM’s low-rank factorization computes 2nd-order crosses in $O(nk)$ vs. $O(n^2)$ for brute-force enumeration, while the Deep component captures higher-order interactions.
- **HGCN** [4]: The MCC merchant category hierarchy (10 L1 categories → 30 L2 subcategories → 109 leaf codes) is inherently tree-structured. Hyperbolic space (Poincaré ball, 8D) embeds trees with exponentially less distortion than Euclidean space [31] — critical for a 550K-node merchant hierarchy. HGCN receives `merchant_hierarchy` features (27D Poincaré embeddings) via FeatureRouter, not product co-holding features (those belong to LightGCN’s `product_hierarchy` group).

⁴The flexibility runs in both directions: in the operational deployment, one task whose label coverage fell below the viability threshold was removed by config-level label-viability gating, leaving 15 tasks in production — likewise a configuration change, not a code change.

Expert	Inductive Bias	Captures
DeepFM [14]	Feature interaction	2nd-order cross features
Temporal Ensemble	Multi-scale temporal	Short/long/disrupted series
HGCN [4]	Hyperbolic hierarchy	MCC merchant category tree (Poincaré)
PersLay [2]	Topological persistence	Behavioral shape patterns
LightGCN [16]	Graph convolution	Collaborative filtering
Causal [45]	DAG constraint	Causal direction between features
Optimal Transport [7]	Distribution matching	Segment distribution shifts

Table 3: Seven heterogeneous experts with distinct inductive biases.

HGCN and LightGCN serve distinct roles that were inadvertently conflated in early experiments. HGCN embeds the MCC merchant category hierarchy (L1: 10 categories $\rightarrow$ L2: 30 subcategories) into Poincaré disk space, capturing tree-structured category relationships. LightGCN learns customer-product bipartite affinity via collaborative filtering on the 24-product co-holding graph. The `feature_groups.yaml` `target_experts` declarations enforce this separation: `merchant_hierarchy` $\rightarrow$ HGCN only, `product_hierarchy` $\rightarrow$ LightGCN only.

- **PersLay** [2]: Topological Data Analysis captures *shape* features (connected components H_0 , cycles H_1 , voids H_2) of customer spending patterns. H_1 loops reveal consumption cycles; H_2 voids detect systematic spending avoidance. These features are provably stable under noise (Stability Theorem).
- **Temporal Ensemble**: Financial time series combine regular snapshots, bursty transactions, and multi-month dormancy. Mamba [13] handles long-range trends in $O(n)$; LNN [15] adapts its time constants to irregular intervals; Transformer captures short-range attention patterns.
- **Causal** [45]: A learnable adjacency matrix $W \in \mathbb{R}^{d \times d}$ (where $d = 32$ compressed feature dimensions) encodes causal direction between features. During training, the NOTEARS acyclicity constraint ($\text{tr}(e^{W \circ W}) - d = 0$, Taylor-approximated to 10th order) is added as a regularization loss, guiding W toward a DAG structure. At inference time, the learned W is frozen and applied as $\hat{z} = z + z \cdot (W \circ W)$, injecting discovered causal structure into the representation without re-running the optimization. This enables “A causes B” explanations rather than “A correlates with B” [34]. (The companion loss-dynamics paper reports empirical findings on this expert, including a silent $W \rightarrow 0$ collapse under the default initialisation and a two-

part patch — NOTEARS reconstruction loss plus initialisation rescale — that restores DAG learning; a follow-up W-amplification experiment grows $\|W\|_F$ $14\times$ at zero primary-task cost. Paper 3 Findings 8 and 11.)

- **LightGCN** [16]: Collaborative filtering via neighborhood aggregation on the customer-product bipartite graph. This expert is stripped to its essentials — no feature transformation and no activation — which has been shown to outperform more complex GCN variants for recommendation.
- **Optimal Transport** [7]: Sinkhorn-regularized Wasserstein distance measures distributional shift between customer spending profiles and segment prototypes, respecting the metric structure of the feature space (unlike KL divergence).

The rationale for heterogeneous experts is also rooted in a hardware constraint: with a single desktop GPU (12GB VRAM), we cannot scale a homogeneous MLP expert to sufficient width/depth for high expressiveness. Instead, each expert leverages a *structural inductive bias* to capture patterns that would require orders of magnitude more MLP parameters. For example, HGCN embeds the MCC merchant category hierarchy in 27 hyperbolic dimensions — achieving in $O(d)$ parameters what Euclidean embeddings require $O(2^d)$ dimensions to represent without distortion [4].

The total parameter count across all seven shared experts is **2.70M** at the V1 349-feature schema (measured). At the V2 1211-feature schema (lag tensor 800D + rolling stats 20D + multihot 24D+31D), the seven routed shared experts total **0.97M**, against 2.96M for the same basket with routing disabled (every expert receiving the full 1211D input) — a 67% reduction from replacing broad shared inputs with expert-specific subsets via FeatureRouter (instantiated and measured at the per-expert dimensions of Table 4). The diversity of learned repre-

sentations is fundamentally richer despite the smaller parameter budget.

3.4.1 FeatureRouter: Heterogeneous Input $\times$ Heterogeneous Architecture

FeatureRouter implements a stronger form of expert specialization: not only does each expert use a **different architecture** (heterogeneous basket), but each also receives a **different subset of input features** (heterogeneous input). This “heterogeneous architecture $\times$

heterogeneous input” design maximizes the signal-to-noise ratio for each expert by eliminating features that carry no information relevant to that expert’s inductive bias.

Feature group assignments are declared in `feature_groups.yaml` via the `target_experts` field, ensuring zero dataset-specific hardcoding in model code. The resulting per-expert input dimensions are summarized in Table 4:

Expert	Input Dim	Feature Groups Routed
DeepFM	977D	• demographics • products • txn_behavior • derived_temporal • gmm • model_derived • txn_lag_tensor • txn_rolling_stats • nba_label_multihot • mcc_top30_multihot
Temporal Ensemble	116D	• txn_behavior • hmm • mamba • model_derived
HGCN	58D	merchant_hierarchy, mcc_top30_multihot
PersLay	32D	tda_global, tda_local
Causal	129D	• demographics • products • txn_behavior • derived_temporal • product_hierarchy • gmm • txn_rolling_stats
LightGCN	955D	• product_hierarchy • graph_collaborative • txn_lag_tensor • nba_label_multihot • mcc_top30_multihot
Optimal Transport	95D	• demographics • products • txn_behavior • derived_temporal • gmm • txn_rolling_stats

Table 4: Per-expert input dimensions after FeatureRouter for the seven shared experts. Total feature space: 1211D (Phase 0 v3/v4, 17 groups). The per-task MLP tower (57D over demographics + products + gmm) is excluded from this shared-expert table.

The sum of per-expert dimensions (2419D) exceeds 1211D because several feature groups are shared across multiple experts where complementary inductive biases benefit from the same signal (e.g., state features are useful to both DeepFM for interaction modeling and Causal for DAG structure inference).

3.4.2 CGC Gate Design: Softmax vs. Sigmoid

Standard PLE uses softmax CGC gates, which impose a competitive constraint: gate weights sum to 1, forcing experts to compete for influence. With homogeneous MLP experts this is reasonable — experts provide redundant representations and the gate selects the best one.

With heterogeneous experts, this competition is harmful: each expert provides unique, non-redundant information (temporal $\neq$ hierarchical $\neq$ topological), and suppressing one means losing irreplaceable signal.⁵ Recent theoretical work [30] proves that sigmoid gating achieves higher sample efficiency by eliminating inter-expert competition.

We implement both gate types for empirical comparison:

- **Softmax:** $w_k = \text{softmax}(W \cdot h)_k$ — competitive, sum-to-one.
- **Sigmoid:** $w_k = \sigma(w_k \cdot h) / \sum_j \sigma(w_j \cdot h)$ — independent evaluation, normalized.

Note: Our sigmoid variant includes post-hoc normalization for comparability with softmax. Nguyen et al.’s theoretical analysis applies to unnormalized sigmoid; the normalized variant used here partially reintroduces competition through the denominator.

The sigmoid variant evaluates each expert’s relevance independently before normalization, allowing multiple experts to receive high weight simultaneously. The theoretical expectation is that sigmoid benefits heterogeneous experts by preserving non-redundant signals. However, our structure ablation (Section 5.4) reveals a surprising reversal: softmax **outperforms** sigmoid in the 12-task heterogeneous setting, because its competitive allocation **protects** minority-type tasks (multiclass, regression) from gradient corruption by the majority type (binary). This finding challenges the theoretical prediction and highlights that task-type heterogeneity, not just expert heterogeneity, determines the optimal gate design.

3.4.3 Temporal Ensemble: Expert-within-Expert

Financial time series exhibit a uniquely complex structure rarely seen in other domains. A single customer may

simultaneously have: monthly balance snapshots (regular, long-range), daily transaction sequences (irregular, bursty), and multi-month dormancy gaps (disrupted). No single temporal model handles all three well:

- **Mamba** [13] (State Space Model): Captures long-range dependencies with $O(n)$ efficiency via selective state spaces. Ideal for monthly/quarterly trends spanning 12+ months.
- **LNN** [15] (Liquid Neural Network): Adaptive time constants that naturally handle irregular sampling intervals and dormancy gaps. When a customer is inactive for 3 months then suddenly active, LNN’s continuous-time dynamics adapt without requiring imputation.
- **Transformer:** Attention-based short-range context extraction. Captures patterns within the most recent 30–90 days of transaction sequences where positional relationships matter.

The three models’ outputs are concatenated and projected, mirroring the heterogeneous expert philosophy at a finer granularity. This design is validated in the temporal component ablation, where removing any single temporal component degrades performance on different task groups.

3.5 Multi-disciplinary Feature Engineering

A distinguishing aspect of this work is the systematic application of methodologies from diverse academic disciplines to financial customer behavior. With the expert basket in place, the rationale for the feature design becomes clear: each discipline’s features are designed to feed a *specific expert* whose inductive bias can best exploit that signal type. Rather than relying solely on standard statistical features (mean, variance, trend), we apply domain-specific mathematical tools that each extract a structurally different signal from the same underlying data (Table 5).

Note: Phase 0 v3/v4 produces 1211 total features across 17 groups (up from 349 in earlier versions). FeatureRouter routes feature subsets to each shared expert (per-expert dims: DeepFM 977D, Temporal 116D, HGCN 58D, PersLay 32D, Causal 129D, LightGCN 955D, OT 95D); model parameters across the seven shared experts: 2.70M (V1, 349D schema, measured); V2 1211D: 0.97M routed vs. 2.96M unrouted (measured). Expert routing is built from `feature_groups.yaml` `target_experts` declarations (group-level), ensuring each expert receives only its designated feature groups.

⁵An a-posteriori canonical-correlation analysis of the trained shared experts supports this premise empirically: the mean pairwise canonical correlation across the 21 expert pairs is 0.20, with most pairs in the LOW band and PersLay and Optimal-Transport near-orthogonal to every other expert. The heterogeneous experts thus occupy largely distinct representation subspaces rather than redundant ones (baseline gate, single seed, benchmark data).

Discipline	Method	Dim	Financial Interpretation
Topology	Persistent Homology (TDA)	32	Behavioral shape persistence: which consumption patterns are transient vs. structural
Hyperbolic Geometry	HGCN embedding	27	MCC merchant category distance: preserving hierarchy in low dimensions
Control Theory	Mamba (State Space)	50	Long-range behavioral dependencies: how past habits influence present
Stochastic Processes	HMM state transitions	25	Latent lifecycle stages: dormant $\rightarrow$ growing $\rightarrow$ mature $\rightarrow$ at-risk
Chemical Kinetics	Reaction rate modeling	6	Spending activation rate, half-life, dormancy reactivation catalysis
Epidemiology	SIR compartment model	5	Product adoption as “infection”: susceptible $\rightarrow$ adopted $\rightarrow$ churned
Criminology	Routine Activity Theory	5	Transaction regularity: burstiness, circadian variance, routine break-points
Signal Processing	FFT + Hilbert transform	8	Spending periodicity: spectral entropy, harmonic power, phase locking
Economics	Friedman Permanent Income	8	Income decomposition: permanent vs. transitory income, consumption smoothing
Graph Theory	LightGCN	66	Collaborative filtering: similar customer behavioral transfer
Statistics	GMM soft clustering	22	Probabilistic segmentation: multi-modal customer distribution

Table 5: Multi-disciplinary feature engineering across 11 academic disciplines.

Several of these applications are, to our knowledge, novel in financial recommendation:

Chemical kinetics models spending behavior as a reaction system: the “activation energy” represents the threshold for a dormant customer to resume spending, while “catalysts” (e.g., payroll deposits) lower this threshold. The half-life of spending intensity after a peak captures decay dynamics that simple moving averages cannot express.

SIR epidemic model treats product adoption as a contagion process: customers are “susceptible” (not yet adopted), “infected” (recently adopted), or “recovered” (churned). The basic reproduction number R_0 of a product category across a customer’s social graph predicts adoption velocity.

Friedman’s Permanent Income Hypothesis [12] decomposes observed income into permanent (stable, long-term) and transitory (bonus, irregular) components using HP filter or Kalman filter. This distinction is critical for financial recommendation: a customer with high transitory income should not be recommended long-term investment products that assume stable cash flow. The FD-TVS scoring system (detailed in companion paper) uses this decomposition to weight recommendations by income stability type.

These features serve a dual purpose beyond predictive contribution: they enrich the *recommendation context* that no standard feature can provide. Internally, “activation energy decrease” (chemical kinetics) signals rising re-engagement likelihood, and “growing SIR infection ratio” (epidemiology) signals accelerating product adop-

tion — but the customer never sees these scientific terms. An `interpretation_registry` maps each scientific feature to business-interpretable language, and the LLM agent generates natural-language reasons accordingly (e.g., “Your recent activity pattern shows signs of recovery”). The scientific frameworks ground the model’s internal reasoning; the customer receives business language.

3.6 Financial DNA Task Grouping

We organize 13 prediction tasks into four groups based on financial customer DNA. Each task group corresponds to one DNA axis (Table 1); within each group, different modality experts (Table 2) contribute differently, and the CGC gate learns the optimal mixture per task:

The four task groups — Engagement, Lifecycle, Value, Consumption — and their constituent tasks are defined in Table 1.

adaTT enforces differentiated transfer: strong intra-group transfer (same DNA perspective) and weaker inter-group transfer (different perspectives, minimizing negative transfer).

Loss-level transfer — terminology and provenance. Our inter-task transfer mechanism differs from the original adaTT [22] on **two** dimensions simultaneously:

1. **Level of operation.** [22]’s adaTT performs representation-level fusion: per-task fusion layers adaptively combine task-specific and shared expert **outputs** through learned gating networks. Our implementation, in contrast, operates at the **loss level**:

$$\mathcal{L}_i^{\text{trans}} = \mathcal{L}_i + \lambda \sum_{j \neq i} w_{i \rightarrow j} \cdot \mathcal{L}_j$$

2. **Affinity source.** [22] learns the fusion weights through differentiable gating on expert activations. We compute $w_{i \rightarrow j}$ from **gradient cosine similarity** between per-task loss gradients, smoothed by EMA — an approach closer to Task Affinity Groupings [10] and GradNorm-style adaptive loss balancing [6] than to [22].

In other words, what we call “adaTT” throughout this paper is not a re-implementation of [22], but a **TAG + GradNorm hybrid for heterogeneous-expert PLE**. We retain the “adaTT” shorthand for continuity with prior internal documentation; readers familiar with [22] should treat the two mechanisms as distinct. The base task losses $\mathcal{L}_i$ are further weighted by learned uncertainty [20], and imbalanced binary tasks (e.g., `churn_signal`) use focal loss [23] with task-specific α and γ .

This departure from [22] was motivated by two heterogeneous-expert considerations: (1) representation-level transfer requires matching hidden dimensions across experts whose native outputs differ in shape and semantics (on-prem H-GCN produces 128D, the remainder 64D⁶), making cross-expert fusion architecturally cumbersome; (2) loss-level transfer is architecturally orthogonal to the PLE layer and can be toggled without modifying the expert routing; gradient-cosine affinity is additionally interpretable (“tasks that help each other move gradients in the same direction”), which matters for regulated MRM documentation.

Note. After correcting five implementation bugs in the on-prem port (gradient-extraction frequency, adaTT config loader, `freeze_epoch` wiring, uncertainty-weighting ordering, and `warmup_epochs` default), our ablation (Section 5.4) shows that this loss-level variant has a **null effect** at the 12-task scale: AUC changes by only -0.001 with adaTT vs. without, well within single-seed measurement noise. An earlier draft reported a -0.019 AUC degradation; that result turned out to be driven by

the five implementation bugs rather than by any intrinsic scalability problem of loss-level transfer. We preserve this note for transparency about the investigative process: it is easy for a paper to publish an “algorithmic finding” that is in fact an implementation artefact, and we believe reporting the correction is more useful than retrofitting a cleaner narrative. GradSurgery (Section 5.4) was tested as a gradient-level alternative but likewise shows no gain over the PLE-only baseline and is not adopted.

3.7 Logit Transfer

While adaTT handles symmetric inter-task relationships, logit transfer captures *directional* dependencies that reflect the natural sequence of customer experience:

These two transfers reflect the natural sequence of customer experience: churn risk gates product-holding stability (lifecycle axis), and next-category intent informs next-best-action selection (consumption axis). Earlier drafts listed an `engagement_score` $\rightarrow$ `nba_primary` edge; this task was removed in 2026-04-12 as a deterministic feature transformation (see Section 3 task reduction), and the `has_nba` $\rightarrow$ `nba_primary` edge was folded into `nba_primary` directly. The transfer directions are not learned from data but specified based on domain knowledge of the customer journey. This is a deliberate design choice: while the *strength* of transfer is learned (via adaTT affinity), the *direction* is fixed by business logic, providing an additional layer of interpretability.

3.8 Inherent Explainability

A central claim of this work is that the heterogeneous expert structure provides *structural interpretability at the analytical-perspective level* — gate weights decompose each prediction into named analytical viewpoints (temporal trends, product hierarchy, causal pathways, etc.) as a natural byproduct of the forward pass. This is not causal explanation in the Pearl sense: most experts operate at Level 1 (association) of the causal ladder. The exception is the Causal expert, which learns a NOTEARS DAG and thus approaches Level 2 (intervention) reasoning. When the Causal expert’s gate weight is

Source $\rightarrow$ Target	Method	Customer Experience	Causal Direction
<code>churn_signal</code> $\rightarrow$ <code>product_stability</code>	<code>output_concat</code>	Retention risk $\rightarrow$ product-holding stability	Causal
<code>next_mcc</code> $\rightarrow$ <code>nba_primary</code>	<code>hidden_concat</code>	Next-category intent $\rightarrow$ next-best-action selection	Sequential / feature sharing

Table 6: Logit transfer relationships reflecting natural customer experience flow (Santander, v12 benchmark — source: `configs/datasets/santander.yaml::task_relationships`).

⁶In the on-prem deployment the H-GCN module was later converted to an embedding-only component — 128D embeddings consumed without a per-task prediction head; the dimensional-heterogeneity argument applies unchanged.

high for a prediction, it indicates not merely correlation but that *structural causal relationships* contributed to the recommendation — a qualitatively stronger form of explanation, valid only for checkpoints whose learned W passes health checks ($\|W\|_F$ in range, active edges present) at evaluation time (see the W -observability limitation and Section 6.4). How sharply this perspective-level decomposition materializes, however, is data- and task-dependent: on operational data the routing is more uniform, so crisp attribution to 1–2 experts is limited to tasks with a dominant expert, while most tasks default to multi-factor explanation (Section 6.4).

Mechanism. The CGC (Customized Gate Control) module computes an attention over the K experts for each task t . With softmax gating: $w_t = \text{softmax}(W_t \cdot h + b_t) \in \mathbb{R}^K$; with sigmoid gating: $w_{t,k} = \sigma(w_{t,k} \cdot h) / \sum_j \sigma(w_{t,j} \cdot h)$. The weight $w_{t,k}$ directly indicates how much expert k contributed to task t ’s prediction.

Business interpretability. Because each expert k has a named inductive bias (e.g., k = “Temporal” means time-series patterns, k = “HGCN” means product hierarchy), $w_{t,k}$ carries business meaning without additional interpretation:

Example: For customer c , the recommendation of a commuter check card is driven by:

Temporal (0.35) — card spending has been increasing over 3 months

HGCN (0.28) — merchant category hierarchy shows concentration in transit and convenience MCCs

DeepFM (0.22) — commute-related product $\times$ channel interaction pattern

Others (0.15)

This example is an idealized profile for tasks with a dominant expert. On operational data (Section 6.4), routing for most tasks is more uniform (normalized entropy 0.67–0.97), so such sharp 3–4-expert decompositions appear only for a subset of tasks.

What gate weights do and do not explain. Gate weights answer “which analytical perspective mattered most” — not “why this customer will buy this product” in a causal sense. For most experts (Temporal, HGCN, LightGCN, DeepFM, PersLay, OT), the weight indicates associative relevance. For the Causal expert specifically, a high weight signals that a learned DAG structure contributed to the prediction, providing a stronger explanatory basis. This is an honest middle ground: richer than SHAP’s feature-level attribution, but short of full counterfactual reasoning.

Comparison with post-hoc methods. SHAP and LIME operate at the *feature* level (“feature_237 con-

tributed 0.12”), requiring an additional mapping step to translate feature attributions into business narratives. Moreover, LIME’s local linear approximation is demonstrably unstable — small input perturbations can dramatically change the explanation [36]. In contrast, gate weights are deterministic functions of the input and change smoothly with input variation.

Dual role of features. An important design insight is that features serve two distinct purposes: (1) predictive contribution to AUC, and (2) explanation material for recommendation reasons. Even features with marginal predictive contribution (e.g., TDA topological features may add only $\Delta\text{AUC} \approx 0.01$) provide irreplaceable context for recommendation reasoning. Internally, TDA persistence signals that a customer’s spending pattern has a stable topological shape — this is then reverse-mapped via the `interpretation_registry` to business language (e.g., “You maintain a stable transaction pattern”) that the customer actually sees. The scientific feature enriches the recommendation context; the customer receives an interpretable description. This dual role motivates our multi-disciplinary feature engineering and is validated in the ablation study.

4 Training Pipeline

4.1 Config-Driven Design

The entire pipeline is controlled by a split-config of three YAML files: `pipeline.yaml` (common model, training, distillation, and AWS settings), `feature_groups.yaml` (feature generation and expert routing), and `configs/datasets/{name}.yaml` (dataset-specific tasks, labels, and ablation scenarios). The split keeps cross-dataset settings in `pipeline.yaml` while letting a new dataset be onboarded by copying `configs/datasets/example.yaml` and redefining task/label schemas only. Adding a new task or expert requires only configuration changes.

4.2 Data Processing

Following the principle of *pandas-free* data engineering for scalability, the pipeline uses a tiered backend strategy: cuDF (GPU columnar) for generator fitting/transformation, DuckDB [35] (CPU columnar with disk spill) for SQL-based aggregation and I/O, and PyArrow for zero-copy parquet loading into training tensors. *pandas* is used only as a last-resort fallback for datasets under 10K rows.

This backend choice has practical implications: DuckDB handles the 1M-row dataset with a 4GB memory limit and automatically spills to disk, enabling processing on commodity hardware. cuDF accelerates generator fitting

(GMM, HMM, Mamba) on GPU, reducing Phase 0 feature engineering from 30 minutes (pandas) to 8 minutes.

4.3 Data Leakage Prevention

Data leakage is a pervasive risk in multi-stage pipelines where features and labels are processed in separate stages. We implement five structural safeguards:

1. **Generator label exclusion:** Feature generators (GMM, model-derived, etc.) automatically exclude all label columns from their input. Without this guard, we observed XGBoost AUC of 1.0 on the generated features — the generator had encoded label information directly into features.
2. **Temporal split with gap:** A configurable `gap_days` parameter ensures no temporal overlap between training and validation windows. For cross-sectional (single-snapshot) data, the system auto-detects and falls back to random splitting.
3. **3-stage normalization with post-normalization group-range rebuild:** Stage 1 detects power-law columns (skew + kurtosis screen followed by log-log R^2) and appends a `log1p` copy of each as an `_log` column; Stage 2 fits `StandardScaler` on the training split only and applies it to continuous (non-binary, non-categorical-ID, non-probability) columns, so validation and test splits are transformed with training statistics; Stage 3 preserves the `_log` offspring at raw magnitude rather than re-scaling, so `FeatureRouter` receives both centered and tail-preserving views of heavy-tailed features. Because Stage 1 appends `_log` columns at the end of the feature matrix, the pre-normalization `feature_group_ranges` become stale; the runner therefore rebuilds each group’s range against the post-normalization column order using a longest-contiguous-block heuristic, which closes a silent routing bug where `FeatureRouter` would otherwise slice the wrong columns for experts whose groups contained power-law features.
4. **LeakageValidator:** A pre-training check computes Pearson correlation between every feature and every label (on a 50K subsample for efficiency). Correlations above 0.95 trigger a warning and investigation.
5. **Forward-transition label criterion:** The fifth safeguard addresses a leakage class that none of the above can catch — *label-definition-internal leakage*, where the label is itself a function of data inside the feature window (e.g. a churn label defined as a snapshot of “no transactions in the past 90 days,” computed over the same months the features describe). Because the leak lives in the label definition rather than in the data flow, no temporal split, gap, or correlation screen blocks it. The formal criterion is that the

label observation window must open strictly *after* the feature cutoff (a *forward transition*); a corollary is that immature target months — those whose label window has not yet closed — must be explicitly rejected rather than silently filled, a maturity constraint that follows from correct label design rather than from any pipeline defect. The benchmark’s task reduction (18 → 13) by removal of deterministic label transformations (Table 1) is the synthetic-data manifestation of the same class, and the criterion was independently found to matter on operational data, where snapshot-defined labels pass every split-based check.

5 Experiments

5.1 Benchmark Data

We construct a 1M-customer synthetic benchmark using a four-layer generative model, inspired by the Synthetic Data Vault framework [32] and conditional generative approaches [42], but with a novel *variance budget* mechanism for controllable difficulty:

1. **Latent Personas:** 6 GMM-fitted personas with 5D continuous latent vector (70% persona-conditioned, 30% independent noise).
2. **Gaussian Copula Demographics:** Correlated demographic variables preserving realistic joint distributions [32].
3. **Vectorized Transactions:** Per-customer transaction sequences as LIST columns.
4. **Variance Budget Labels:** Each label’s predictability is controlled via $\text{logit} = \sqrt{f_{\text{obs}}} \cdot z_{\text{obs}} + \sqrt{f_{\text{lat}}} \cdot z_{\text{lat}} + \sqrt{f_{\text{noise}}} \cdot \varepsilon$, with post-hoc label noise flipping.

The benchmark underwent several iterations (v2 through v14; see Appendix F for version history). The final generation schema, v12, introduces Financial DNA axis-aligned situation variables — each customer receives an independent situation per DNA axis (engagement: steady/surging/declining/volatile; lifecycle: stable/growing/consolidating/transitioning; value: stable/ascending/descending/shock; consumption: consistent/exploring/focusing/switching). Situations modulate transaction sequence patterns (temporal spacing, amount trends, MCC distribution shifts, product acquisition rates) without altering static demographic features, creating signals recoverable only through Phase 0 generators (Mamba, TDA, HGCN, GMM). Label interactions use 3rd–5th order continuous multiplicative terms with `obs_frac=0.15`, producing meaningful class distributions for all generated tasks. This is a design guarantee of the synthetic benchmark and does not transfer automatically to real data: the default state of operational labels was degeneration and extreme imbalance (single-value labels,

Tier	Labels	f_{obs}	f_{noise}	XGB AUC
Easy	segment ⁷	determ.	n/a	0.95–1.0
Core	churn_signal	0.04	0.68	0.58–0.65
Hard	will_acquire_*	0.03	0.72	0.50–0.56
V.Hard	next_mcc, top_mcc_shift	0.02	0.78	0.50–0.51

Table 7: Variance budget per label tier (benchmark v4 per-tier allocation). XGB AUC ceiling validates difficulty control. The final generation schema (v12), shared by all results in this paper, uses a unified global allocation (obs_frac=0.15, lat_frac=0.35, noise_frac=0.50) with 3rd–5th order multiplicative label interactions, producing meaningful class distributions across all generated tasks (13 at the time; 12 after segment_prediction was excluded in the v13 revision — described in the paragraph below).

missing class coverage — Section 6.4), so label-viability gating (class count, coverage, and variance checks) is a mandatory pre-training stage.

Version terminology. v12 froze the data generation schema; later revisions did not change how the benchmark data is generated. v13 and v14 are *phase0 preprocessing* revisions applied to v12-generated data: v13 is the first such revision, and v14 additionally introduces HMM mode-split fitting, a GMM component reduction to $K = 14$ (removing dead clusters), and the exclusion of probability-valued columns from standard scaling. Result tables are therefore labelled with the preprocessing revision on which they were measured (e.g. “v14 phase0” in Table 8, Table 10, and Table 11; the structure ablation of Table 9 predates the v14 fixes), while “v12” elsewhere refers to the generation schema common to all of them.

Label signals combine linear observable features (obs_frac=0.15) with latent components (lat_frac=0.35) recoverable through behavioral sequence patterns. Observable interactions are 3rd to 5th order continuous multiplicative (e.g., income $\times$ spend $\times$ tenure $\times$ products $\times$ stability for investment propensity), requiring deep nonlinear capacity that tree-based models and shallow networks cannot efficiently approximate. The remaining variance (noise_frac=0.50) is irreducible noise, setting a hard upper bound on model performance.

5.2 Experimental Setup

The ablation validates whether each expert provides a distinct “why” for different task types — confirming that heterogeneous experts are not merely a performance trick but a structural requirement for multi-faceted persuasion.

- **Data:** 1M customers, 1211 features (1210 after the v14 preprocessing revision; Phase 0 over the v12 generation schema, 17 groups; preprocessing revision per table caption), 12 tasks (7 binary + 2 multiclass + 3 regression; segment_prediction was excluded in the v13 revision for leakage risk).
- **Hardware:** NVIDIA RTX 4070 (12GB VRAM, 64GB RAM) local.
- **Training:** 10 epochs, batch 5632, lr 0.0005, AMP (FP16), warmup 3 epochs (cosine annealing), no early stopping. adaTT scenarios use warmup=3 epochs, grad_interval=10. GradSurgery scenarios use warmup=2 epochs, conflict_threshold=0.0.⁸
- **Loss weighting:** Uncertainty weighting (Kendall et al.) with per-task loss weights applied on top of learned precision — matching the on-premise reference formula.
- **Metrics:** Selected to match each task’s production semantics. Binary classification uses AUC (threshold-independent and imbalance-robust). The multiclass tasks (nba_primary with 7 product groups, next_mcc with top-50 merchant categories) use F1-macro for classification quality and NDCG@K with top-K accuracy [18] for ranking quality — both tasks play a dual classification/recommendation role, and each table caption states which metric it reports. Regression uses MAE. All metrics are reported per task type; a global average across all 12 tasks is avoided because metrics have incompatible semantics across types.

5.3 Joint Feature + Expert Ablation (RQ1 + RQ2)

Bottom-up ablation adds one expert (with its matched feature generator) to the DeepFM baseline; top-down

⁷segment_prediction is a 4-class customer lifecycle label derived from behavioral clustering. It is part of the benchmark generation schema and was used as a task through the v12 experiments, but the strong feature-label correlation its XGB ceiling (0.95–1.0) implies led to its exclusion from the final training task set in the v13 revision (the segmentation signal is instead carried by nba_primary’s 7-class distribution); it is not part of this paper’s 12-task results.

⁸Reproducibility note: scheduler and optimizer hyperparameters resolve with a fixed three-tier precedence — explicit hyperparameter arguments override YAML configuration, which overrides library defaults — so the values reported here are the values actually applied.

Scenario	Avg AUC	Avg F1m	Avg MAE	Val Loss
<i>Baselines</i>				
Full (7 experts, PLE softmax)	0.8200	0.1777	0.3161	18.23
DeepFM only (all generated features)	0.8076	0.1798	0.3167	21.57
Raw inputs only (no generated features)	0.6894	0.1725	0.4299	21.82
<i>Bottom-up: DeepFM + single expert (sorted by AUC desc)</i>				
DeepFM + HGCN	0.8001	0.1826	0.3198	26.08
DeepFM + Temporal	0.7999	0.1813	0.3223	23.41
DeepFM + PersLay (TDA)	0.7991	0.1845	0.3236	26.89
DeepFM + LightGCN	0.7979	0.1905	0.3208	29.29
DeepFM + OT	0.7949	0.1923	0.3266	31.97
DeepFM + Causal	0.7932	0.1875	0.3243	29.92
<i>Top-down: Full minus one expert (sorted by ΔAUC asc)</i>				
Full – HGCN	0.7965	0.1941	0.3241	33.41
Full – OT	0.7935	0.1921	0.3291	37.46
Full – TDA	0.7923	0.1967	0.3296	35.01
Full – Temporal	0.7919	0.1916	0.3229	32.47
Full – LightGCN	0.7908	0.1888	0.3307	35.16
Full – Causal	0.7862	0.1976	0.3287	37.49

Table 8: Joint feature + expert ablation results on Santander v14 phase0 (1M rows, 17 months, HMM mode-split + GMM K=14 + probability scaler-excluded). Bold = best in column. The full 7-expert ensemble strictly dominates every single-expert and 6-expert configuration on Avg AUC and Val Loss, demonstrating that expert synergy is a non-trivial property of the design rather than an averaging artefact.

ablation removes one expert from the full 7-expert model. Avg AUC averages 7 binary tasks; Avg F1m averages 2 multiclass tasks (nba_primary, next_mcc); Avg MAE averages 3 regression tasks (cross_sell_count, mcc_diversity_trend, product_stability). All scenarios use PLE softmax gating with adaTT disabled, batch 2048, 10 main epochs (3 warmup + 7 phase2), AMP enabled, on RTX 4070.

Expert contribution analysis (Table 8) reveals five patterns.

Pattern 1: Generated features carry the bulk of the signal; expert ensembling adds the final layer. The raw-input baseline (no generated features, only demographics + product holdings + lag/rolling tensors + multi-hot labels) reaches AUC 0.6894; adding the generator suite (HMM, GMM, TDA, Mamba, model-derived, etc.) and a single DeepFM expert lifts AUC to 0.8076 (+0.118 absolute); expanding to the full 7-expert heterogeneous basket adds another +0.0124 to reach 0.8200. The decomposition (raw $\rightarrow$ generated features: +0.118; single $\rightarrow$ ensemble: +0.012) confirms two architectural claims simultaneously: generator-driven feature engineering carries the dominant lift, and the heterogeneous ensemble extracts a non-trivial residual that no single expert recovers.

Pattern 2: The full 7-expert ensemble dominates Avg AUC; partial pools never recover the gap. Full (0.8200) beats DeepFM-only (0.8076) by +0.0124 AUC — a meaningful gap that no bottom-up scenario closes. Counter-intuitively, every single-expert addition to DeepFM *decreases* Avg AUC below the DeepFM-only baseline (best single addition: +HGCN at 0.8001, -0.0075 relative to baseline); and every single-expert removal from Full also stays below the full-pool number (best post-removal: $-$ HGCN at 0.7965, -0.0235 relative to Full). This convex shape indicates the seven experts contribute *interaction signal*, not just additive features: pairwise mixtures inject feature noise that the gate cannot down-weight without seeing all seven together.

Pattern 3: Causal expert is the most load-bearing component (RQ2 update). Removing the Causal expert costs -0.0338 AUC (Full $\rightarrow$ Full–Causal: 0.8200 $\rightarrow$ 0.7862), the largest top-down degradation among all six experts. The ranking of removal impact (Causal $-0.034 >$ LightGCN $-0.029 >$ Temporal $-0.028 >$ TDA $-0.028 >$ OT $-0.026 >$ HGCN -0.024) identifies Causal/LightGCN as the most essential and HGCN as the most redundant for aggregate AUC — a different ordering from the earlier on-prem prototype run’s LightGCN-dominant narrative, attributable to the v14 phase0 fixes

that exposed valid causal DAG signal previously masked by HMM mode collapse and probability-column scaling artefacts.

Pattern 4: F1-macro inverts AUC — minority-class tasks prefer ablated configurations. The full ensemble has the *worst* Avg F1m (0.1777) of all 14 successful scenarios; every ablated configuration scores higher F1m, with peak F1m at Full—Causal (0.1976) and DeepFM+OT (0.1923). This inversion signals that softmax gating over-averages multiclass logits at the 7-expert pool size: adding more experts smears the per-class signal that minority multiclass labels (nba_primary, next_mcc) need. This is the same softmax-redistribution mechanism documented in the on-prem benchmark, manifesting here as a binary–multiclass trade-off rather than absolute degradation.

Pattern 5: Validation loss separates ensemble configurations more than AUC. Val Loss ranges 18.23 (Full) → 37.49 (Full—Causal) — a 2× spread, versus only 4% spread in Avg AUC. The val_loss gap suggests the ensemble’s benefit is concentrated in the long tail of low-confidence predictions that AUC’s rank-only metric flattens. This is consistent with the production setting where calibrated probability estimates drive downstream LGBM distillation and FD-TVS scoring, and with the SR 11-7 model risk requirement that ensembles demonstrate measurable Loss-level (not just AUC-level) improvement to justify added complexity.

5.4 Task × Structure Cross Ablation (RQ3)

Nine structure variants are compared across three dimensions: gate type (shared-bottom / softmax / sigmoid), inter-task transfer (none / adaTT loss-level / GradSurgery gradient-level), and uncertainty weighting (all variants use the corrected implementation where per-task loss weight is applied on top of learned precision). All variants use the full 7 heterogeneous expert basket with 10 epochs, batch size 5632, and AMP enabled. Avg AUC covers binary tasks only; NDCG@3 is reported for nba_primary (7-class product recommendation); Avg F1m covers multiclass tasks; Avg MAE covers regression tasks. Results appear in Table 9.

Three findings emerge from the structure ablation.

Finding 1: PLE softmax outperforms sigmoid in heterogeneous MTL. This reverses the conventional wisdom from homogeneous-task PLE literature [41], where sigmoid gates are preferred for their non-competitive expert weighting. In our 12-task setting with 7 binary + 2 multiclass + 3 regression tasks, softmax’s competitive selection *protects* minority-type tasks (multiclass) by isolating expert assignments, preventing binary-task gradients from corrupting multiclass-relevant experts. Crucially, the uncertainty weights converge to identical values (nba_primary: 0.335) under both gate types, confirming that the improvement is purely structural.

Variant	Avg AUC	Best NDCG@3	Final NDCG@3	Avg F1m	Avg MAE
<i>Gate type comparison (no inter-task transfer)</i>					
Shared Bottom	0.6711	0.7014	0.6831	0.2007	0.9602
PLE Softmax	0.6729	0.7144	0.6814	0.2009	0.9598
PLE Sigmoid	0.6728	0.7131	0.6820	0.2021	0.9601
<i>adaTT loss-level transfer (corrected)†</i>					
SB + adaTT	0.6708	0.7048	0.6791	0.1999	0.9606
Sigmoid + adaTT	0.6717	0.7127	0.6646	0.2013	0.9598
Softmax + adaTT	0.6722	0.7022	0.6856	0.2015	0.9597
<i>GradSurgery gradient-level projection‡</i>					
SB + GradSurgery (exp.)	0.6704	0.6963	0.6860	0.1986	0.9614
Softmax + GradSurgery (exp.)	0.6726	0.6918	0.6830	0.2027	0.9588
Sigmoid + GradSurgery (exp.)	0.6721	0.6888	0.6811	0.2020	0.9603

Table 9: Structure ablation across gate type and inter-task transfer variants. Bold = best per metric within gate comparison. †The adaTT rows reflect **corrected** numbers after five implementation bugs inherited from the on-prem port were identified and fixed (gradient-extraction frequency, adaTT config loader path, `freeze_epoch` wiring, uncertainty-weighting ordering, and `warmup_epochs` default). An earlier draft reported a pre-fix Sigmoid+adaTT AUC of 0.6541 and a −0.019 AUC degradation; the corrected values are presented here. ‡GradSurgery uses batch 4096 (vs. 5632) due to retain_graph VRAM overhead.

Finding 2: adaTT loss-level transfer has a null effect at 12-task scale. After correcting the five porting bugs listed in the table caption, adding adaTT changes AUC by only -0.001 (sigmoid+adaTT 0.6717 vs. sigmoid-alone 0.6728; softmax+adaTT 0.6722 vs. softmax-alone 0.6729) and does not shift any other metric beyond single-seed measurement noise. The pre-correction -0.019 AUC drop reported in an earlier draft is now attributed to these implementation bugs rather than to an algorithmic property of loss-level transfer, which was the previously hypothesised cause. Our loss-level gradient-cosine variant therefore neither helps nor meaningfully hurts at 12-task scale: the auxiliary loss weighting produces transfer signal that is **indistinguishable from zero** relative to PLE-only baselines. This should not be generalised as a negative result for representation-level adaTT [22], which we do not evaluate; that mechanism learns fusion weights end-to-end on expert activations rather than from gradient cosine statistics and is a distinct algorithm (see Section 3.6). A direct comparison is deferred to follow-up work.

Finding 3: Correcting uncertainty weighting yields larger gains than architecture changes. The single largest performance improvement ($+0.018$ NDCG@3, $+0.031$ F1-macro) came not from gate type or transfer mechanism, but from fixing a subtle implementation gap where per-task loss weights were silently ignored when uncertainty weighting was active. This finding underscores that *loss balancing correctness* can dominate *architectural sophistication* in heterogeneous MTL — a practical lesson often overlooked in architecture-focused papers.

5.5 Expert Contribution Analysis: Beneficial vs. Negative Transfer (RQ4)

We assess each expert’s contribution by $\Delta\text{AUC} = \text{AUC}(\text{Full}) - \text{AUC}(\text{Full}-\text{Expert})$ on the v14 phase0 run (Full AUC = 0.8200, joint_full 10-epoch). Positive ΔAUC means removing the expert hurts AUC, i.e. the expert is *beneficial* for binary tasks; zero or negative ΔAUC would indicate dispensability or negative transfer. Results are summarized in Table 10.

Removed Expert	Avg AUC after	ΔAUC	Interpretation
–Causal	0.7862	+0.0338	most load-bearing
–LightGCN	0.7908	+0.0292	collaborative-filtering signal essential
–Temporal	0.7919	+0.0281	Mamba+LNN fusion non-trivial
–TDA	0.7923	+0.0277	topology persists across tasks
–OT	0.7935	+0.0265	Sinkhorn alignment helpful
–HGCN	0.7965	+0.0235	least essential but still positive

Table 10: Expert contribution analysis on Santander v14 phase0: ΔAUC relative to full 7-expert model (AUC = 0.8200). Every expert contributes positively ($\Delta\text{AUC} > 0$); none is dispensable. Range 0.024–0.034 is more than an order of magnitude larger than the earlier on-prem prototype run’s $|\Delta\text{AUC}| < 0.001$ noise floor, attributable to v14 fixes (HMM mode-split, GMM K=14, probability scaler-excluded) that exposed real expert differentiation.

All six ΔAUC values are positive (+0.024 to +0.034), confirming that *every* heterogeneous expert provides beneficial signal at the 7-expert pool size. The ranking has three diagnostic implications.

Diagnostic 1: Causal and LightGCN dominate. Causal (+0.034) and LightGCN (+0.029) together account for *sim*50% of the ensemble’s improvement margin over DeepFM-only. The Causal expert’s NOTEARS DAG over MCC categories captures sequential causation (e.g., grocery *to* gas *to* restaurant patterns) that DeepFM’s pairwise factorization cannot represent; LightGCN’s bipartite user–MCC graph encodes collaborative similarity invisible to per-row featurization. The earlier on-prem prototype run reported a LightGCN-dominant ranking with magnitudes < 0.001 ; the v14 ranking inverts the dominant expert (Causal $>$ LightGCN) because the v14 phase0 fixes corrected three signal-masking artefacts: (1) HMM journey/lifecycle mode collapse that fed redundant inputs to Temporal expert, (2) GMM K=20 with 7 dead clusters that fed near-zero probability columns to Causal/OT, and (3) probability

columns z-scored to $pm1.7$ instead of stayed in $[0, 1]$ for downstream gate weighting.

Diagnostic 2: HGCN is least essential but still beneficial. HGCN’s $+0.0235$ ΔAUC is the smallest, suggesting product-hierarchy encoding contributes less aggregate signal than MCC-level graph or causal structure. This is consistent with Santander’s flat product taxonomy (24 leaf SKUs, 2-level hierarchy) limiting HGCN’s hyperbolic curvature advantage. The same expert is expected to dominate on production data with a 4–5 level product tree.

Diagnostic 3: F1-macro vs. AUC trade-off is real, not an artefact. Table 8 shows that all top-down ablated configurations beat the full ensemble on F1-macro (0.1976 vs. 0.1777). This is *not* a contradiction with the positive ΔAUC ranking: the full ensemble’s softmax-gate redistribution diffuses minority-class logits across more experts, lowering F1-macro on imbalanced multi-class tasks (nba_primary class 0 dominance, next_mcc long-tail). The architectural implication is that production deployment should select the metric appropriate for the downstream task — for FD-TVS scoring driven by binary purchase-intent AUC, the full ensemble wins; for next-best-action ranking driven by multiclass coverage, a 6-expert variant (Full–Causal) may be preferable, especially when the secondary metric is rank-aware (NDCG, top-k recall).

5.6 Epoch-Budget Sensitivity (RQ4 follow-up)

The local 10-epoch runs (the Local 10ep column of Table 11) suggested an unexpected ordering: shared_bottom (Avg AUC 0.8139) outperformed PLE softmax (0.7980) on aggregate AUC, contradicting the v13 prototype’s PLE-favoured narrative. Three competing hypotheses were considered: **(A) signal-cleaning made gate-routing redundant**, **(B) PLE+CGC gates need more than 10 epochs to converge**, and **(C) the gate parameters overfit faster than they help**. A SageMaker re-run of four core scenarios (shared_bottom, PLE-softmax, PLE-full = sigmoid+CGC+GTE+LT+HMM-projectors, PLE-full+adaTT) at both 10 and 15 epochs on ml.g4dn.2xlarge — identical hardware, identical fp32 DuckDB-streaming pipeline (Section 4.2 memory budget) — yields Table 11. The matched 10-epoch column on g4dn.2xlarge isolates the epoch-budget effect from any precision-path drift between the local fp64 pq.read_table baseline and the SageMaker fp32 pipeline.

Variant	Local 10ep AUC	SM 10ep AUC	SM 15ep AUC	ΔAUC (15ep - 10ep)	SM val_loss
shared_bottom	0.8139	0.8197	0.8015	−0.0182	17.64
PLE softmax	0.7980	0.8233	0.8249	+0.0016	17.09
PLE full (sigmoid + 5 toggles)	—	0.8216	0.8203	−0.0013	17.43
PLE full + adaTT	0.8200	0.8223	0.8213	−0.0010	17.60

Table 11: Epoch-budget sensitivity, v14 phase0 (1M rows). *Local 10ep*: RTX 4070 fp64 pq.read_table. *SM 10ep / 15ep*: ml.g4dn.2xlarge fp32 DuckDB-stream. ΔAUC compares matched 15–10 on SageMaker. All four PLE variants plateau between 10 and 15 epochs ($|\Delta AUC| < 0.002$), while shared_bottom drops -0.0182 — the architecture-class signature of this dataset.

The matched-pair comparison rejects **Hypotheses A and B**. At 10 epochs on SageMaker, PLE softmax (0.8233) already exceeds shared_bottom (0.8197) by $+0.0036$ AUC, so signal-cleaning did not make gate-routing redundant (A); and the local-10ep ordering was an artefact of the fp64 pq.read_table path, not an epoch-convergence deficit of PLE (B). What Table 11 does reveal is **Hypothesis C** (overfitting asymmetry): shared_bottom drops -0.0182 AUC moving $10 \rightarrow 15$ epochs with val_loss climbing $16.58 \rightarrow 17.64$, while PLE softmax stays flat ($+0.0016$ AUC, val_loss $16.04 \rightarrow 17.09$). Three implications follow.

Implication 1: shared_bottom requires early stopping or regularization at the 10-epoch budget. The -0.018 AUC drop between 10 and 15 epochs and the matching val_loss climb ($16.58 \rightarrow 17.64$) indicate that the unfettered shared representation memorises training noise faster than it generalises. A regularization-aware deployment of the shared-bottom baseline would pin its checkpoint at the early-stopping epoch.

Implication 2: adaTT loss-level transfer remains null. PLE-full (+ adaTT) at 0.8213 sits within $pm0.001$ of PLE-full without adaTT (0.8203). This reproduces the v13 finding (Finding 2 of Table 9) at the longer 15-epoch budget on cleaner v14 data: the loss-level gradient-cosine

signal is indistinguishable from zero relative to the PLE-only baseline.

Implication 3: Auxiliary toggles (GTE + LT + HMM-projectors) do not improve PLE softmax on aggregate AUC. PLE-full (sigmoid + CGC + GTE + LT + HMM-proj) at 0.8203 is -0.0046 below PLE-softmax-alone at 0.8249. Adding GTE only (PLE-sigmoid-GTE: 0.8214) and GTE+LT (PLE-sigmoid-GTE-LT: 0.8204) traces a flat line slightly below the softmax baseline. The auxiliary stack contributes complexity without AUC benefit on this data, so production should ship the simpler PLE-softmax configuration and defer the routing-aux machinery to per-task fine-tune.

A regularized variant **PLE-softmax-reg** (dropout 0.3, weight_decay 1×10^{-4}) further improves to AUC 0.8256 and val_loss 16.57 (vs. PLE-softmax 0.8249, val_loss 17.09). This $+0.0007$ AUC and -0.52 val_loss confirms the gate-overfitting hypothesis: even at its best architecture, PLE leaves calibrated-loss room on the table that conventional regularization recovers without architectural addition.

adaTT loss-level transfer monotonically increases val_loss across gate types (Table 12). Adding adaTT raises val_loss by $+0.44$ on shared_bottom (17.64 $\rightarrow$ 18.08), $+1.54$ on PLE-sigmoid (16.99 $\rightarrow$ 18.53), and $+1.93$ on PLE-softmax (17.09 $\rightarrow$ 19.02), while moving Avg AUC by at most $pm0.005$ in any direction. The auxiliary loss-level signal therefore destabilises the calibration tail without delivering aggregate-AUC compensation, and the v13 finding (Finding 2 of Table 9) that adaTT is null on AUC at this scale is now upgraded to **adaTT is val-loss harmful at this scale** under the longer training budget.

Base	Vanilla AUC	+ adaTT AUC	Vanilla val_loss	+ adaTT val_loss
shared_bottom	0.8015	0.8097	17.64	18.08
PLE sigmoid	0.8240	0.8217	16.99	18.53
PLE softmax	0.8249	0.8197	17.09	19.02
PLE full (sigmoid+5)	0.8203	0.8213	17.43	17.60

Table 12: adaTT pairwise comparison at 15 epochs on SageMaker g4dn.2xlarge. Bold entries mark the better choice within each row (lower val_loss, higher AUC). adaTT delivers AUC gains only on shared_bottom and PLE-full, but increases val_loss in every pairing — most severely on PLE-softmax ($+1.93$) and PLE-sigmoid ($+1.54$).

5.7 Explainability Analysis (RQ5)

The CGC gate produces interpretable routing weights that enable direct attribution of “which expert contributed how much” per task. Under the recommended softmax configuration, weights sum to one, making relative contributions directly comparable. We examine per-task gate weight distributions across all tasks to identify (a) which experts dominate which task types, and (b) whether the learned routing aligns with domain intuition. This alignment is task- and data-dependent, however: in operational observations, merchant_affinity and channel weight the Temporal expert highly and balance_util weights the Causal expert, loosely matching intuition, but a specific mapping such as “the temporal expert for churn prediction” does not hold (Section 6.4).

Gate weights from the CGC extraction layers provide per-task expert utilization profiles. Tasks with low entropy ratio (e.g., top_mcc_shift at 0.347) concentrate on 1–2 experts, providing clear attribution: the recommendation is driven primarily by the dominant expert’s feature group. Tasks with high entropy (e.g., will_acquire_payments at 0.882) leverage diverse experts, requiring multi-factor explanation. This entropy-based explainability directly maps to the Financial DNA task groups and enables the rule-based fallback (Layer 3) to select appropriate explanation templates per task. On operational data, however, routing is more uniform, so the set of tasks for which such crisp single-expert attribution holds is limited (Section 6.4).

Detailed per-task gate weight profiles and dominant expert assignments are provided in the Gate Entropy

Task	Layer 1 $H / \log K$	Layer 2 $H / \log K$	Interpretation
top_mcc_shift	0.347	0.542	Single expert dominance
product_stability	0.428	0.705	Concentrated routing
segment_prediction	0.612	0.332	Layer 2 concentrates
nba_primary	0.877	0.724	Diverse expert usage
will_acquire_payments	0.882	0.688	Diverse expert usage
CGC Attention (all tasks)	1.000	1.000	Uniform (undifferentiated)

Table 13: Per-task CGC gate entropy ratios ($H_t / \log K$; higher = more uniform). Measured on the v12 phase0 13-task checkpoint (segment_prediction was excluded from the task set from v13 onward).

Analysis (RQ6) below, which jointly addresses both routing collapse and task-expert specialization.

5.8 Gate Entropy Analysis (RQ6: Does routing collapse occur?)

Beyond function space collapse (prevented structurally by heterogeneous experts), routing collapse can occur when CGC gate weights concentrate on a single expert. We measure per-task gate entropy:

$$H_t = - \sum_k w_{t,k} \log w_{t,k}$$

where $w_{t,k}$ is the gate weight for expert k on task t . Maximum entropy $\log K$ ($K=7$ experts, $\log 7 \approx 1.95$) indicates uniform expert utilization; entropy near zero indicates routing collapse. We report normalized entropy ratio $\frac{H_t}{\log K} \in [0, 1]$ to facilitate comparison across different numbers of experts.

5.8.1 CGC Layer Gate Specialization

Measured on the PLE-softmax checkpoint (v12 phase0, i.e. the 13-task configuration prior to the exclusion of segment_prediction; later revisions shrink only the task set to 12, with the gate structure unchanged), CGC layer gate weights show meaningful task-level specialization with entropy ratios ranging from 0.33 to 0.88 across the 13 tasks and 2 CGC layers (Table 13). This confirms that routing collapse has not occurred: all tasks engage multiple experts, while distinct tasks show distinct routing patterns. Re-measured on the operational teacher ($K=6$, 12 tasks), the normalized entropy is 0.67–0.97 — more uniform than synthetic, with 7 of 12 tasks exceeding 0.90; the no-collapse conclusion holds on both, but crisp dominant-expert specialization is confined to a few tasks (Section 6.4). Operational deployment also demonstrated two failure modes distinct from routing collapse: (i) zeroed regressions, where a modality input wiring regression pins an expert’s input to the zero vector and its output goes constant (observed historically for Temporal, PersLay, and HGCN), and (ii) degeneration of an expert’s

internal parameters (the Causal W global collapse, Section 3.4). The heterogeneous design’s “structural collapse guarantee” therefore holds only conditional on continuous verification of input wiring integrity and monitoring of expert-internal parameter health. Low-entropy tasks concentrate on 1–2 experts; high-entropy tasks draw broadly across the expert basket. CGC attention weights remain perfectly uniform across all tasks, indicating that expert selection operates at the extraction layer, not at the attention aggregation layer.

The divergence between extraction-layer and attention-layer routing is structurally informative. *Low-entropy tasks* (top_mcc_shift: 0.347, product_stability: 0.428, segment_prediction layer 2: 0.332) concentrate predominantly on 1–2 experts, suggesting that these tasks are solved by a narrow analytical lens (e.g., temporal trend or graph structure alone). *High-entropy tasks* (will_acquire_payments: 0.882, nba_primary: 0.877) draw broadly from the expert basket, consistent with the nature of next-best-action and payment acquisition tasks that integrate multiple behavioral signals simultaneously.

Notably, CGC *attention* weights remain perfectly uniform (entropy ratio = 1.000 across all tasks and both layers), indicating that the attention aggregation mechanism has not differentiated — functionally equivalent to simple averaging of expert outputs. This suggests the model relies on the extraction-layer gates to perform meaningful routing, while the attention mechanism acts as a residual blending layer. Future work may benefit from regularizing or removing the attention aggregation to reduce parameter redundancy.

6 Discussion

6.1 Findings Summary

The ablation experiments yield the following principal findings. The five claims below summarise the system-level outcomes; readers seeking the full hypothesis

discrimination across A–E (signal cleaning, epoch budget, ensemble, gate overfitting, regularisation), per-task v14 phase0 numbers across all 12 tasks, the 9-way fusion-mechanism comparison (BRP-detached, NEAS, M1 complement, ECEB, AdaTT-sp), and the 30-epoch loss-metric decoupling diagnostic should consult the companion Paper 3 (Findings 1–7) where each is reported with the full ablation table and per-task breakdown.

Loss balancing correctness dominates architecture choice. The single largest performance improvement across all experiments came from correcting a subtle implementation gap in uncertainty weighting (Kendall et al.), where per-task loss weights from the pipeline configuration were silently ignored. The corrected formula — $\text{loss_weight} \times (\text{precision} \times L + \log_var)$ with $\log_var$ clamped to $[-4, 4]$ and precision clamped to $[10^{-3}, 100]$ — matches the on-premise reference implementation and yielded +0.018 NDCG@3 and +0.031 F1-macro on the shared-bottom baseline alone. This underscores that in heterogeneous MTL with mixed loss types, *getting the loss weighting right* matters more than architectural sophistication.

Softmax gating outperforms sigmoid in heterogeneous MTL. This finding contradicts the conventional preference for sigmoid gating in PLE literature [41] and recent sigmoid MoE work [30]. In our 12-task setting (7 binary + 2 multiclass + 3 regression), softmax’s competitive expert selection *protects* multiclass tasks by assigning dedicated experts, while sigmoid’s cooperative blending allows binary-task gradients to corrupt multiclass-relevant representations. The effect is metric-specific: softmax leads in NDCG@3 (+0.013 over shared-bottom) while sigmoid and softmax are near-identical in AUC. We attribute the reversal to *task-type heterogeneity*: when all tasks share the same loss type (as in prior work), cooperative blending is beneficial; when loss types differ fundamentally, competitive isolation is protective.

Loss-level transfer (adaTT) provides no measurable benefit at 12-task scale. After correcting the five porting bugs in the caption of Table Table 9, the AUC change with adaTT vs. without is -0.001 across both gate types (sigmoid and softmax) — within single-seed measurement noise. The mechanism — adding weighted auxiliary losses $L_i^{\text{adaTT}} = L_i + \lambda \sum_{j \neq i} w_{i \rightarrow j} L_j$ with $w_{i \rightarrow j}$ estimated from gradient cosine similarity — produces transfer signal that is **not distinguishable from zero** relative to PLE-only baselines. An earlier draft reported a -0.019 AUC drop and attributed it to a 132 task-pair affinity estimation limit; with the bug-fixed implementation, the effect vanishes, so the attribution is no longer supported by our data. The honest conclu-

sion is operational: at 12-task heterogeneous scale, the gradient-cosine loss-level variant (TAG [10] + GradNorm [6] in spirit) neither helps nor hurts, adds compute cost, and can be disabled without loss. The representation-level adaTT of [22] is a distinct mechanism that learns fusion weights end-to-end on expert activations and is not directly evaluated here; we leave that comparison to follow-up work. A 9-way comparison that places loss-level adaTT alongside seven other fusion-augmentation mechanisms (AdaTT-sp, M1 complement, ECEB, BRP-MV, BRP-detached, NEAS, NEAS+BRP-detached) is reported in companion Paper 3, Finding 7; the v14 verdict isolates only two non-trivial recipes (M1 complement, NEAS) on disjoint axes and confirms the null verdict for loss-level adaTT.

GradSurgery: gradient-level projection as an alternative. To address the loss-level transfer limitation, we tested task-type gradient surgery (GradSurgery), which replaces loss-level transfer with gradient-level conflict resolution. Instead of estimating 132 pair-wise affinities, GradSurgery aggregates task gradients into 3 type-groups (binary, multiclass, regression) and applies PCGrad-style projection only between conflicting groups. This reduces the estimation problem from $O(k^2)$ to $O(3^2) = O(1)$ relative to the number of tasks. Softmax+GradSurgery achieves AUC 0.6726 (vs. softmax-alone 0.6729) and the highest F1-macro (0.2027) and lowest MAE (0.9588) across all variants, while best NDCG@3 (0.6918) is lower than softmax-alone (0.7144). GradSurgery does not degrade performance either, but — like the corrected adaTT — it shows no meaningful advantage over PLE softmax alone. The retained computation graph overhead constrains batch size on 12GB GPUs, and the NDCG gap suggests that gradient-level projection may not fully capture the ranking signal that PLE’s representation-level gate routing provides. GradSurgery was therefore not adopted for production. The most robust configuration remains PLE softmax without transfer mechanisms.

Operational motivation validated. The shared-bottom baseline (AUC 0.6711, NDCG@3 0.7014) already exceeds per-task XGBoost ceilings (on the synthetic benchmark), validating the operational motivation for MTL consolidation: a single model replaces 12 individual models with no performance penalty, enabling unified training, serving, and monitoring. PLE softmax provides incremental gains (NDCG@3 +0.013) at no additional serving cost. This superiority relationship did not transfer to the operational distribution, however: in a clean OOT measurement on operational real data, raw-feature hard-label LGBM outperformed the teacher on 10 of 11 active tasks (Section 6.4). The synthetic “exceeds

XGB ceilings” result should be read as a product of the label-generation design, which embeds deep nonlinear interactions; the verified operational value of MTL consolidation lies not in per-task predictive superiority but in unified training, representation, and explanation plus two-tier serving — the teacher anchoring distillation, explanation, and monitoring while tree models serve as the adopted students.

6.2 Training Dynamics: Loss–Metric Decoupling and Epoch Selection

A 30-epoch training run (cosine LR with $T_0 = 10$, warm restarts) reveals a systematic decoupling between validation loss and discriminative metrics that has practical implications for best-model selection in heterogeneous MTL.

Note: This 30-epoch diagnostic run is separate from the 10-epoch ablation results reported in Tables 8–9. It was conducted to analyze convergence dynamics and does not affect the reported ablation numbers.

Observation: metrics peak at epoch 10, not epoch 30. Validation loss decreases monotonically across all 30 epochs ($32.11 \rightarrow 22.68$, -29.4%), suggesting continued learning by a standard stopping criterion. However, discriminative metrics peak at epoch 10 and subsequently degrade: AUC peaks at epoch 10 (0.6726) and falls to 0.6687 at epoch 30 (-0.4pp), while NDCG@3 peaks at epoch 10 (0.6976) and drops to 0.6673 at epoch 30 (-3.0pp).

Mechanism: regression task dominance of the loss function. The aggregate validation loss is dominated by the three regression tasks, which produce large MAE values relative to binary cross-entropy magnitudes. Continued training genuinely improves regression MAE ($0.9596 \rightarrow 0.9580$ at epoch 30), but at the cost of overfitting to synthetic benchmark noise in the classification and ranking tasks. The benchmark’s formula-based feature–label relationships are fully learned by epoch 10; subsequent epochs overfit to residual noise that has no counterpart in the discriminative metrics.

Cosine LR restarts amplify the effect. The cosine schedule with $T_0 = 10$ causes sharp metric oscillations at cycle boundaries (epochs 11, 20, 29): NDCG@3 drops abruptly at each restart as the learning rate spikes and disrupts the representation alignment that the gate weights have accumulated. This suggests that cosine warm restarts, while beneficial for escaping loss plateaus,

may periodically damage ranking-relevant representations in the heterogeneous setting.

Implication: composite metric for best-model selection. Validation loss alone is an unreliable stopping criterion when task types are heterogeneous and regression loss dominates. We recommend selecting the best checkpoint using a composite metric — e.g., $\alpha \cdot \text{AUC} + \beta \cdot \text{NDCG} + \gamma \cdot (1 - \text{MAE})$ — rather than minimum `val_loss`. In our setting, this corresponds to selecting epoch 10 over epoch 30, recovering $+0.4\text{pp}$ AUC and $+3.0\text{pp}$ NDCG@3. Adopting a composite metric also requires defining a fallback rule for missing constituent metrics: in an operational run, every component of the composite happened to depend on one task’s (nba) metrics, and when that label went missing the composite degenerated to $-\infty$, silently falling back to the previous best checkpoint (Section 6.4). Explicit weight renormalization or promotion of a next-priority metric must be defined for the missing-component case.

6.3 Practical Implications

Resource-constrained development. This system was built without dedicated ML infrastructure budget, on a single desktop GPU, by a three-person team augmented with AI development agents. This demonstrates that complex multi-task recommendation systems are no longer exclusive to organizations with large ML teams and GPU clusters. The key enablers were: (1) config-driven architecture minimizing code changes, (2) AI agents handling parallel implementation tasks under human architectural guidance, (3) heterogeneous expert design achieving expressiveness through structural bias rather than parameter scale, and (4) knowledge distillation eliminating GPU serving costs.

Defense over offense in production ML. In financial recommendation services, maintaining AUC above an operational threshold matters far more than pushing it higher. The heterogeneous expert design directly supports this *defensive* posture: expert contribution diagnostics (ablation reveals which inductive biases are beneficial vs.

harmful per data regime), drift detection triggering automatic retraining, and the champion-challenger gate (offline `ModelCompetition` significance test with a two-tier fidelity gate — Tier 1 blocking on operational metrics scoped to the distilled tasks, including the student’s own ground-truth calibration error; Tier 2 reporting cross-architecture teacher–student diagnostics without blocking — plus an operator `--force-promote` override, every decision written to an HMAC-signed audit chain) that

⁹The two-tier split responds to a structural property of cross-architecture distillation: PLE→LGBM teacher–student agreement plateaus at roughly 0.75–0.82, so a single blocking gate that includes cross-architecture fidelity fails every candidate and trains operators to reach

governs whether a newly trained model may replace the current production champion.⁹ The audit chain itself is tamper-evident beyond signature checking: verification re-computes every entry’s HMAC in constant time (so wholesale re-chaining after an edit is detected, not just broken links), well-known development keys are rejected outright in production and staging environments, and the persisted log is provisioned with WORM object-lock storage. Compliance hooks along the serving path are wired in by default; the stricter fail-closed enforcement — refusing to serve when the kill-switch state or audit trail is unavailable — ships as an opt-in configuration recommended for production. The architecture prioritizes *not getting worse over getting better* — a perspective that aligns with financial regulators’ emphasis on model risk management and operational resilience. Specifically, the kill-switch and fail-closed enforcement correspond to the FSC guideline’s Auxiliary-means principle (§3, emergency stop for high-risk AI) and Financial-stability principle (§5.2, backup model plus post-hoc emergency stop) [11]; as noted above, the fail-closed enforcement is currently opt-in, with default wiring a planned follow-up.

Feature engineering philosophy. A deeper lesson from this work is that *what to observe* matters more than *how to model*. The heterogeneous expert architecture is effective not because of architectural novelty alone, but because each expert is paired with features that ask a **specific question** about the customer: “Is their spending decaying with a half-life?” (chemical kinetics), “Is product adoption spreading like contagion?” (SIR epidemiology), “Is their income permanent or transitory?” (Friedman PIH [12]). The model merely combines answers to these questions. If the questions are limited to conventional statistics — mean, variance, trend — even the most sophisticated architecture cannot capture the structural complexity of customer behavior. This suggests that multi-disciplinary feature engineering, often treated as preprocessing, deserves elevation to a first-class architectural design decision, co-designed with the expert types that will consume each feature group.

Generalizability beyond finance. While this work targets financial product recommendation, the underlying principle — understanding a single user from multiple perspectives simultaneously via heterogeneous experts — applies to any domain where personalization requires multi-faceted user modeling. Healthcare (diagnosis + treatment + risk + lifestyle), education (knowledge level

+ learning style + engagement + career path), and insurance (risk assessment + product fit + claim prediction + retention) all face structurally similar challenges: multiple interdependent prediction tasks about the same individual, where different analytical lenses capture different aspects. We expect the heterogeneous expert PLE pattern to generalize to these domains, with domain-specific expert types replacing the financial experts while retaining the structural benefits of collapse resistance and inherent explainability.

Lessons from design iteration. The final architecture emerged through repeated trial and rejection. The first candidate — Black-Litterman Bayesian model combination — was abandoned because the blended posterior made individual model contributions opaque, failing the explainability requirement that drives the entire project. The second candidate — N-model ensemble — was rejected for $N \times$ management overhead. These failures led to the key reframing: combine experts *inside* a single model, not *outside* it.

An early ablation (before the uncertainty weighting fix) showed sigmoid outperforming softmax, which appeared to confirm the NeurIPS 2024 finding [30] that softmax competition harms heterogeneous experts. However, after correcting the loss weighting implementation, softmax reversed to outperform sigmoid on the primary ranking metric (NDCG@3), demonstrating that **loss balancing bugs can masquerade as architectural findings** — a cautionary lesson for MTL practitioners.

A particularly instructive failure: a configuration bug that disabled PLE collapsed the 7-expert basket into a single MLP, making all ablation scenarios in that early run produce identical AUC (0.913).¹⁰ This was only discovered through systematic result comparison — reinforcing the principle that ablation results must be verified against expected variation before drawing conclusions.

Infrastructure choice: extending GPU ROI beyond LLM serving. Financial institutions have increasingly invested in GPU infrastructure, but predominantly for generative AI workloads — chatbots, document summarization, and LLM-based customer service. Recommendation model training remains largely CPU-based (XGBoost, LightGBM), leaving GPU capacity underutilized outside of LLM inference.

for the override, defeating the control. Promotion thresholds themselves are unchanged. See the companion Paper 2 Discussion “Fidelity Gate Semantics under Cross-Architecture Distillation.”

¹⁰The 0.913 AUC reflects an earlier benchmark version (v4) with a different scenario count and data distribution; the final benchmark (v12 generation schema with v13/v14 phase0 preprocessing) reports 23 scenarios with AUC in the 0.67–0.82 range depending on preprocessing revision.

Our architecture demonstrates that the same GPU hardware already deployed for LLM serving can simultaneously support 12-task MTL training, requiring only 12GB VRAM on a single consumer-grade GPU (RTX 4070). This extends the return on existing GPU investments without additional procurement — a practical consideration for institutions where new hardware acquisition requires multi-month approval cycles. For institutions without on-premise GPU, our SageMaker + Lambda serverless architecture provides an alternative: GPU Spot instances (60–70% discount) are provisioned on-demand for training and released immediately after, while serving runs on CPU-only Lambda functions via knowledge distillation to LGBM (detailed in companion paper). On the v14 phase0 benchmark the distillation stage itself also completes end-to-end on a CPU instance (m5.4xlarge, 62 minutes, \$0.30 on Spot), with the distilled students tracking the teacher at a mean AUC gap of 0.58pp (down from a 2.6pp mean at v13) — GPU dependence is confined to teacher training.

Operational simplicity. The config-driven design (a split-config of three YAML files controls the entire pipeline) enables teams with 1–2 ML engineers to operate the system. Adding a new product category, task, or feature group requires only configuration changes, not code modifications. This stands in contrast to model ensemble approaches where each model is a separate management point with its own training pipeline, serving endpoint, and monitoring dashboard.

6.4 Preliminary Operational Teacher (Research-Only)

To check whether the synthetic-benchmark conclusions hold directionally in a real deployment, we performed two operational observations: (i) a **June 2026 preliminary probe** — gate and liveness measurements of the deployed on-premise teacher (full-expert `best_composite` checkpoint); and (ii) a **July 2026 integrated run** — the v9 teacher, re-running the full pipeline from preprocessing through teacher training on a two-month operational window (real contract schema with a 4,035-dimension input, training on the 2026-01 snapshot of 17.8M rows, validation/test on 2026-02, six shared experts, 11 active tasks, phase 1 for 15 + phase 2 for 8 epochs). **Both observations are research comparisons, not promotion candidates** — v9 failed the performance gate and was recorded as candidate-only, with the blocking reason being nba label availability rather than performance (itself live evidence that the manual-approval promotion gate works as designed). They differ from the synthetic benchmark in task composition (task definitions and naming are entirely different, and the

active count itself shifts — 12 in the June probe, 11 in the July run), number of live experts (operational $K=6$ — LightGCN’s signal absorbed into the unified HGCN input — vs. synthetic $K=7$), and label distribution, so **absolute numbers are not directly comparable — only signs and patterns are cited.**

Independent reproduction of no expert collapse (June probe). Forwarding the operational teacher with side-band inputs wired, all six heterogeneous experts retain non-degenerate outputs (output std 0.051–0.333, weakest is HGCN). Per-task gate entropy ratios span 0.67–0.97; no task collapses onto a single expert (Table 14). The core RQ6 conclusion — no collapse — reproduces independently on both synthetic and operational data. However, operational routing is more uniform than synthetic (0.33–0.88): 7 of 12 tasks exceed 0.90, so crisp attribution to 1–2 experts is limited to the few tasks with a dominant expert (spending_category with PersLay, merchant_affinity with Temporal, life_stage with PersLay).

The July integrated run’s OOT measurements (v9). With OOT labels (2026-03) only partially generated, out-of-time evaluation was possible for just 3 of the 11 active tasks: churn AUC 0.9422 (PR-AUC 0.7214), balance_util R^2 0.7360, and merchant_affinity R^2 0.1485. Churn’s calibration error is ECE 0.2485, uncalibrated — calibration is mandatory before any downstream use of the probabilities. Training ran under a per-epoch batch cap of 969 (roughly an 11% subsample of train), leaving undertraining a possibility, and the wall-clock from preprocessing through teacher was about 87 hours (fully reproduced through the Airflow DAG path).

High absolute metrics reflect task easiness, not architectural superiority — now confirmed by a measured tree comparison. Some operational metrics appear high (June probe churn AUC 0.89; July v9 0.9422), but this is attributable to a narrow, skewed operational distribution and extreme label imbalance: the churn positive rate is below 1% (0.54% in the June observation), and the majority class of the multiclass tasks accounts for 96–99%. The large AUC–PR-AUC gap (June: 0.89 vs. 0.47) and high calibration error (ECE 0.25–0.31) corroborate this. The June-era hypothesis that “a simple GBM baseline may already match or exceed” was confirmed by measurement in the July integrated run: in a clean OOT evaluation, raw-feature hard-label per-task LGBM models outperformed the teacher on 10 of 11 active tasks (11/11 when restricted to active segments), and stacking with teacher logits added no gain. The sparsity structure — 93–94% of customers dormant — also produced Simpson reversals, where aggregate and

Task	Top expert	Gate wt.	$H_t/\log K$
spending_category	PersLay	0.590	0.673
merchant_affinity	Temporal	0.534	0.768
life_stage	PersLay	0.427	0.825
balance_util	Causal	0.370	0.862
channel	Temporal	0.329	0.875
retention	DeepFM	0.331	0.901
spending_bucket	DeepFM	0.375	0.917
engagement	HGCN	0.281	0.924
churn	DeepFM	0.338	0.926
nba	DeepFM	0.294	0.931
consumption_cycle	OT	0.259	0.943
timing	PersLay	0.235	0.973

Table 14: Operational teacher (research-only, June 2026 probe, $K=6$, 12 tasks): per-task top expert and gate entropy ratio. All six experts non-degenerate (output std 0.051–0.333). Entropy 0.67–0.97 shows no collapse but is more uniform than synthetic (0.33–0.88). Task names do not map 1:1 to synthetic Table 13, so this is a separate table and absolute numbers are not directly compared.

fine-grained segment verdicts flip; the synthetic benchmark does not reproduce this environmental factor. Indeed, in the operational gate, the generated-feature DeepFM ranks top for most tasks (churn, retention, nba, spending_bucket). This operational observation is therefore liveness evidence — no collapse, all experts active — not a performance-superiority claim; the operational value of MTL consolidation lies in unified training, representation, and explanation plus two-tier serving (Section 6.1). Gate-type comparison (softmax vs. sigmoid), the adaTT ablation, and GradSurgery were not re-run on operational data, so those conclusions remain scoped to the synthetic benchmark.

Label availability governs the task portfolio. The default state of real-data labels was degeneration and extreme imbalance: engagement degenerated to a single distinct value; merchant_affinity had 2 distinct values at 6.1% coverage; nba observed only 5 of 12 classes (majority class 96.1%); life_stage shrank from 6 to 3 classes. Campaign collection ending (2026-03-10) extinguished the ctr/cvr task sources, and retention auto-deactivated as churn’s inverse. What determines the task portfolio is thus label-source availability, not architectural capacity, and label-viability gating (class count, coverage, variance checks) is a mandatory pre-training stage. Distillation-related operational measurements — confidence compression in $T = 5$ soft-label storage, the fidelity gate’s degenerate-pass blind spot (high fidelity achievable by constant imitation), and monotonic clean-OOT degradation as the soft-label weight grows — are detailed in the companion Paper 2’s operational observations section. Access to the operational environment ended with this observation window, so no further operational measurements are possible — the measurements

in this section constitute the final operational evidence for this system.

6.5 Limitations

Synthetic benchmark only. All experiments use a 1M-customer synthetic benchmark with formula-based feature-label relationships. While the benchmark controls difficulty via a variance budget mechanism and validates against XGBoost AUC ceilings, it cannot reproduce the complex cross-feature interactions present in real financial customer data. In particular, the formula-based label generation may understate the benefit of expert specialization — real data where temporal patterns, graph structures, and causal pathways carry genuinely different signals may amplify PLE’s advantage over shared-bottom. A **preliminary** validation on proprietary financial data was performed on the operational teacher (Section 6.4), but that teacher is a candidate excluded from production promotion by label-availability gating; the observation window spans only two months (training 2026-01, validation/test 2026-02), OOT evaluation covers just 3 of 11 active tasks, so generalization across seasonality and regime shifts is unverified; training ran under a per-epoch batch cap of 969 (11% subsample), leaving undertraining a possibility; and extreme label imbalance (churn positive rate 0.54%) demands caution in metric interpretation. Full production validation remains future work. That future validation will target labels defined under the forward-transition criterion of Section 4.3 (label windows opening strictly after the feature cutoff), so that reported operational gains cannot be inflated by label-definition-internal leakage.

Skewed operational distribution and measured tree superiority. The high absolute metrics seen in the

operational observation (churn AUC 0.89/0.9422) stem from a narrow, skewed distribution and extreme label imbalance (majority class 96–99% on multiclass tasks), indicating low task difficulty rather than architectural superiority. Indeed, in the clean OOT measurement, raw-feature hard-label LGBM outperformed the teacher on 10 of 11 active tasks (in the operational gate the generated-feature DeepFM ranks top for most tasks). The synthetic-benchmark result that “MTL exceeds XGBoost ceilings” does not transfer to the skewed operational distribution. Per-expert leave-one-out retraining has not been run, however, so operational evidence for per-expert contribution also remains incomplete.

Label-quality defects contaminate metrics.

Degenerate labels — engagement (distinct=1), merchant_affinity (6.1% coverage), nba (5 of 12 classes observed) — limit the interpretability of those tasks’ metrics. The ctr/cvr family is unreproducible after collection ended (2026-03-10), and retention, as churn’s inverse, shrinks the effective count of independent tasks.

Uncalibrated probabilities. The operational teacher’s churn ECE is 0.2485 (June probe: 0.31), uncalibrated. Calibration is mandatory before any downstream use of probabilities (threshold decisions, expected-value scoring).

W observability gap. The Causal expert’s W -health indicators ($\|W\|_F$, active edges) are absent from the evaluation report and were confirmed only via a separate probe (the v9 W is frozen at $\|W\|_F \approx 0.17$ with 7 edges). Wiring them in as mandatory evaluation-report fields is future work.

Scope of operational reproduction; no absolute comparisons. The gate-type comparison, adaTT, and GradSurgery conclusions were not re-run on operational data (synthetic-only). Absolute-number comparisons between synthetic and real data are artifacts of the difficulty-design difference and are invalid (signs and patterns only).

adaTT evaluation at limited epoch budget. The corrected adaTT implementation’s null effect may be partly attributable to the 10-epoch training budget: with 50+ epochs, affinity estimation may stabilize sufficiently for 132 task pairs and potentially produce a measurable positive effect. The 10-epoch budget reflects practical constraints (ablation throughput on a single GPU) and is sufficient for convergence of non-transfer variants.

Single-GPU training. The current implementation runs on a single GPU (RTX 4070, 12GB). DDP support is architecturally designed but not yet experimentally validated.

LLM dependency. Recommendation reason generation relies on LLM inference, introducing latency and cost trade-offs (detailed in companion paper).

6.6 Ethics and Data Statement

All **quantitative benchmark** experiments in this version use **fully synthetic benchmark data** generated via Gaussian Copula with a fixed random seed (seed=42). The preliminary observations of the operational teacher (Section 6.4 — the June 2026 probe and the July 2026 integrated run) report only aggregate metrics and gate statistics as a research-only comparison, and include no raw customer data or personally identifying information. Full production-promotion validation remains future work. The benchmark data generator and all configuration files are publicly available in the repository. The system is designed for deployment on low-risk check card products only; investment and insurance recommendations are explicitly excluded from the operational scope (see companion paper, Section 6.3).

6.7 Future Work: Scaling Considerations

The current architecture is intentionally lightweight — each expert uses a compact, domain-specific design (e.g., 2-layer MLP for DeepFM deep branch, pre-computed embeddings for HGCN) rather than scaling parameters. This is a deliberate design choice, not merely a resource constraint: structural inductive biases substitute for raw parameter count.

When scaling to larger institutions with dedicated GPU infrastructure, the natural progression is:

- **Data enrichment** — longer transaction histories, broader product coverage, richer interaction signals — before model capacity increase;
- **Per-expert input dimension expansion** — wider feature groups feeding each expert’s specialized inductive bias;
- **Expert depth/width scaling** — deeper layers within individual experts, particularly Temporal (longer sequence modeling) and HGCN (deeper hierarchy encoding);
- **Multi-GPU training** via DDP, which is architecturally supported but not yet validated.

Notably, scaling the expert basket size (adding more expert types) is less promising than scaling individual expert capacity, because the heterogeneous design already covers the major mathematical perspectives relevant to financial behavior.

7 Conclusion

We presented Heterogeneous Expert PLE, a multi-task learning architecture for financial product recommendation that consolidates 12 heterogeneous tasks (7 binary + 2 multiclass + 3 regression) into a single model, replacing 12 individual models while maintaining or exceeding per-task performance (on the synthetic benchmark; for its position under the skewed operational distribution and the realigned value of consolidation, see Section 6.4).

Three contributions emerge from this work. **First**, the heterogeneous expert basket — seven architecturally distinct experts with FeatureRouter-based input specialization — provides structural collapse resistance and inherent explainability through business-interpretable gate weights, without reliance on post-hoc attribution methods. Collapse resistance reproduces independently on both synthetic and operational data (Section 6.4); crisp single-expert attribution, however, holds for the few tasks with a dominant expert, while most tasks require multi-factor explanation.

Second, ablation at 12-task scale reveals that loss-level inter-task transfer (adaTT) has a null effect once five implementation bugs inherited from the on-prem port are corrected (−0.001 AUC change, within single-seed noise), while the single largest improvement comes from *correcting loss balancing* rather than architectural changes. This finding challenges the common assumption that architectural sophistication is the primary lever for MTL improvement, and highlights that implementation correctness of loss weighting deserves the same scrutiny as architecture design.

Third, we tested GradSurgery, a task-type gradient projection method that addresses inter-task interference at the gradient level rather than the loss level, reducing the $O(k^2)$ task-pair estimation problem to $O(1)$ type-group projections. GradSurgery matches PLE-only performance (F1-macro 0.203 vs. the corrected adaTT’s 0.201) but does not surpass PLE softmax alone on the primary ranking metric (NDCG@3 0.692 vs. 0.714) and incurs significant VRAM overhead from the retained computation graph. GradSurgery is therefore not adopted for production. The most effective configuration is therefore PLE softmax *without* transfer mechanisms — structural expert isolation through gating proves more robust than numerical gradient manipulation.

Softmax gating outperforms sigmoid in this heterogeneous setting, reversing findings from homogeneous-task literature — a result we attribute to softmax’s protective

isolation of minority-type tasks (multiclass) from majority-type gradients (binary).

The architecture, benchmark data, and ablation framework are released as open source.¹¹ Two companion papers address the remaining scope. **Paper 2**¹² covers the downstream serving pipeline: knowledge distillation [17] to LGBM [19], the two-tier fidelity gate semantics under cross-architecture distillation, multi-agent recommendation reason generation, regulatory compliance mapping for Korea’s FSC Financial-Sector AI Guideline and EU AI Act requirements, and an HMAC-signed hash-chained per-prediction audit surface, optionally enriched by a causal attribution (CEH) / reliability (CG) pair from the companion loss-dynamics paper. **Paper 3** reports empirical findings from scaling PLE to 13 heterogeneous tasks — loss dynamics and gate selection behaviour (Findings 1–6), a 9-way comparison of fusion augmentations that identifies two non-additive positive recipes (Finding 7), a causal expert reinterpretation arc from dead-parameter diagnosis through CEH attribution, CG guardrail, and W-amplification to counterfactual probing and a negative-result CEH variant (Findings 8–13) whose outputs can optionally enrich Paper 2’s audit surface, and a measurement analysis of the cross-architecture distillation fidelity gate (Finding 14).

Author Contributions

Seonkyu Jeong (PM / Lead Architect / Data Scientist): Conceived the project direction, designed the heterogeneous expert architecture and two-axis decomposition framework, selected multi-disciplinary feature engineering approaches based on structural isomorphism, defined task groups and regulatory compliance mapping, led AI-augmented development methodology, and wrote the manuscript. Overall technical leadership and Scrum-based rapid feedback coordination.

Euncheol Sim: Data ingestion pipeline, feature engineering implementation, feature business reverse-mapping, and vector database pipeline management.

Youngchan Kim: Model training, mathematical verification, and knowledge distillation implementation.

All authors collaborated through Scrum sprints with rapid feedback cycles.

Funding

This research received no external funding, grants, or institutional infrastructure support. All costs — including AI development tools (Claude Code, Gemini, Cursor sub-

¹¹https://github.com/bluethestyle/aws_ple_for_financial

¹²DOI: <https://doi.org/10.5281/zenodo.19622051>

scriptions), hardware peripherals, mobile data connectivity, AWS SageMaker cloud training (Spot instances), S3 storage, and operational expenses — were borne entirely by the first author’s personal funds. Development was conducted on a single desktop-grade GPU (NVIDIA RTX 4070, 12GB VRAM) with no dedicated ML infrastructure budget, no institutional network support, and no institutional cloud computing allocation. Data collection was constrained to a legacy HIVE environment with no access to Spark or Impala, requiring custom parallel query logic to overcome I/O bottlenecks.

Acknowledgments

The authors thank Yeon-Jin Kim for consistently providing valuable insights on industry trends and marketing perspectives that informed the system’s design direction.

The authors express deep gratitude to Euncheol Sim and Youngchan Kim, who dedicated countless nights and weekends to this project with unwavering commitment despite the absence of formal employment or compensation.

Finally, the authors wish to acknowledge the AI tools that made this research possible for a team of three with no institutional support. Gemini served as a tireless brainstorming partner throughout ideation and design direction discussions. Claude Opus and Sonnet were with us through nearly every stage — systematically structuring the architecture, implementing code, diagnosing subtle bugs, and helping articulate ideas into written form. Cursor provided the seamless development environment that kept our workflow free of bottlenecks. What could have remained one person’s unrealized vision was brought to life through the collaboration of a small team and these tools. The architectural decisions, domain knowledge, experimental design, and research direction were led by the human authors; the AI tools amplified what would otherwise have been impossible to achieve at this scale.

Appendix

A. Configuration Schema

The entire system is controlled by a split-config of three YAML configuration files: `pipeline.yaml` (model architecture, training hyperparameters, distillation, AWS deployment settings — the settings shared across datasets), `feature_groups.yaml` (feature group definitions, generator parameters, expert routing rules), and `configs/datasets/{name}.yaml` (dataset-specific task definitions, label derivations, and ablation scenarios). `load_merged_config()` deep-merges the dataset file into

the common settings at run time. Full configuration files are available in the accompanying repository.

B. Ablation Scenario Definitions

The 23 ablation scenarios (14 joint feature+expert + 9 structure cross) are organized into two phases:

Phase 1 — Feature + Expert Joint Ablation (14 scenarios): Full model (7 experts, PLE softmax) and DeepFM-only baseline; 6 bottom-up scenarios (DeepFM + one of Temporal, HGCN, PersLay, LightGCN, Causal, OT with matching features); and 6 top-down scenarios (full minus one of the same 6 experts).

Phase 2 — Task x Structure Cross Ablation (9 scenarios): Nine structural variants — shared-bottom, PLE-softmax, PLE-sigmoid (no inter-task transfer); shared-bottom+adaTT, PLE-softmax+adaTT, PLE-sigmoid+adaTT (loss-level transfer); and three GradSurgery variants (shared-bottom+GradSurgery, softmax+GradSurgery, sigmoid+GradSurgery) — all using the full 12 tasks and 7 heterogeneous experts with 10 epochs.

C. Benchmark Data Generation

The synthetic benchmark uses a four-layer generative model: (1) latent personas via 6-component GMM with 5D continuous latent vector (70% persona-conditioned, 30% independent noise); (2) Gaussian Copula demographics preserving realistic correlations; (3) vectorized transaction sequences as LIST columns; (4) variance-budget labels with post-hoc noise flipping (6–8%). Seed=42 ensures reproducibility. Full generation code is available in the accompanying repository.

D. AMP FP16 NaN Diagnosis and GradScaler Tuning

Heterogeneous experts with on-prem-aligned activation functions (ODE-based LNN, linear HGCN output, Softplus TDA weights) produce wider intermediate value ranges than homogeneous MLPs. Under AMP (FP16), this initially caused GradScaler overflow that cascaded into NaN loss — observed from epoch 2 in PLE configurations with default GradScaler settings. All reported experiments use AMP FP16 training, which provides approximately $2\times$ throughput on the T4/RTX GPU. The NaN issues were resolved by tuning three GradScaler parameters: `init_scale=1024` (lower initial scale reduces overflow at startup), `growth_interval=2000` (slower scale growth prevents rapid re-escalation after overflow events), and `max_scale=4096` (hard ceiling prevents runaway scaling for wide-range expert activations). Additionally, per-task log-variance parameters in the uncertainty weighting head are clamped to $[-4, 4]$ to prevent Softplus overflow from propagating through the

loss. With these mitigations, zero NaN batches are observed across all ablation runs.

E. Structural Isomorphism Verification

To address the concern that cross-disciplinary feature engineering may be “plausible analogy rather than rigorous isomorphism,” we present the governing equations side by side:

Chemical Kinetics -> Spending Activation: Arrhenius: $k = Ae^{-E_a/RT}$ where E_a = activation energy. Spending: $p_{\text{reactivate}} = Ae^{-E_{\text{threshold}}/S}$ where S = stimulus intensity. Both describe the probability of a state transition as a function of an energy barrier and an external driving force.

SIR -> Product Adoption: Epidemiology: $d\frac{I}{dt} = \beta SI - \gamma I$ Adoption: $d\frac{A}{dt} = \beta_{\text{exposure}} \cdot U \cdot A - \gamma_{\text{churn}} \cdot A$ where U = unadopted customers, A = adopted, β = contact/exposure rate. The compartmental dynamics are identical; only the variable names change.

F. Benchmark Version History

The benchmark generation schema underwent four major iterations before converging on v12; two phase0 preprocessing revisions (v13, v14) followed on top of the frozen schema:

- **v2:** Uniform-random MCC codes and fixed transaction amounts across personas, producing near-random labels for MCC-dependent tasks.
- **v3:** Persona-weighted MCC distributions (4–5× boost) with temporal stickiness (30%), persona-dependent transaction amounts, and quantile-based `spend_level` boundaries. Improved MCC task signal but still yielded near-uniform label distributions for acquisition tasks.
- **v4:** Sharpened persona MCC preferences (8–12× boost), increased temporal stickiness to 60%, raised per-product acquisition rates (8–12%), and widened `top_mcc_shift` detection window to 30 transactions.
- **v12** (final generation schema): Financial DNA axis-aligned situation variables, 3rd–5th order continuous multiplicative label interactions with `obs_frac`=0.15, producing meaningful class distributions for all generated tasks (13 at the time). The generation schema is frozen here; v13 and v14 below modify preprocessing only (`segment_prediction` was excluded from the task set in v13, so subsequent experiments use 12 tasks).
- **v13** (phase0 preprocessing): first preprocessing revision over the frozen v12 schema; basis of the structure ablation (Table 9).
- **v14** (phase0 preprocessing): HMM mode-split fitting, GMM $K = 14$ dead-cluster removal, probability columns excluded from standard scaling, plus a

streaming-memory refactor for SageMaker 1M-row training; basis of Table 8, Table 10, and Table 11.

Bibliography

- [1] AI21 Labs. 2024. Jamba: A Hybrid Transformer-Mamba Language Model. *arXiv preprint arXiv:2403.19887* (2024).
- [2] Mathieu Carrière, Frédéric Chazal, Yuichi Ike, Théo Lacombe, Martin Royer, and Yuhei Umeda. 2020. PersLay: A Neural Network Layer for Persistence Diagrams and New Graph Topological Signatures. *Journal of Machine Learning Research* (2020).
- [3] Rich Caruana. 1997. Multitask Learning. *Machine Learning* 28, 1 (1997), 41–75.
- [4] Ines Chami, Rex Ying, Christopher Ré, and Jure Leskovec. 2019. Hyperbolic Graph Convolutional Neural Networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [5] Wei Chen, Yong Liu, and Hao Zhang. 2024. Deep Learning-Powered Financial Product Recommendation System in Banks: Integration of Transformer and Transfer Learning. *Journal of Organizational and End User Computing* 36, 1 (2024).
- [6] Zhao Chen, Vijay Badrinarayanan, Chen-Yu Lee, and Andrew Rabinovich. 2018. GradNorm: Gradient Normalization for Adaptive Loss Balancing in Deep Multitask Networks. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, 2018.
- [7] Marco Cuturi. 2013. Sinkhorn Distances: Light-speed Computation of Optimal Transport. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2013.
- [8] European Parliament and Council. 2024. Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (AI Act).
- [9] William Fedus, Barret Zoph, and Noam Shazeer. 2022. Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity. *Journal of Machine Learning Research* (2022).
- [10] Christopher Fifty, Ehsan Amid, Zhe Zhao, Tianhe Yu, Rohan Anil, and Chelsea Finn. 2021. Efficiently Identifying Task Groupings for Multi-Task Learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- [11] Financial Services Commission (Korea). 2026. Financial Sector AI Guidelines.

- [12] Milton Friedman. 1957. *A Theory of the Consumption Function*. Princeton University Press.
- [13] Albert Gu and Tri Dao. 2024. Mamba: Linear-Time Sequence Modeling with Selective State Spaces. (2024).
- [14] Huifeng Guo, Ruiming Tang, Yunming Ye, Zhenguo Li, and Xiuqiang He. 2017. DeepFM: A Factorization-Machine based Neural Network for CTR Prediction. In *Proceedings of the 26th International Joint Conference on Artificial Intelligence (IJCAI)*, 2017.
- [15] Ramin Hasani, Mathias Lechner, Alexander Amini, and others. 2021. Liquid Time-constant Networks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2021.
- [16] Xiangnan He, Kuan Deng, Xiang Wang, Yan Li, Yongdong Zhang, and Meng Wang. 2020. LightGCN: Simplifying and Powering Graph Convolution Network for Recommendation. In *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2020.
- [17] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 2015. Distilling the Knowledge in a Neural Network. In *NIPS Deep Learning and Representation Learning Workshop*, 2015.
- [18] Kalervo Järvelin and Jaana Kekäläinen. 2002. Cumulated gain-based evaluation of IR techniques. *ACM Transactions on Information Systems* 20, 4 (2002), 422–446.
- [19] Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. 2017. LightGBM: A Highly Efficient Gradient Boosting Decision Tree. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [20] Alex Kendall, Yarin Gal, and Roberto Cipolla. 2018. Multi-Task Learning Using Uncertainty to Weigh Losses for Scene Geometry and Semantics. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018.
- [21] Jiaxin Li and others. 2024. HoME: Hierarchy of Multi-Gate Experts for Multi-Task Learning at Kuaishou. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2024.
- [22] Yankai Li and others. 2023. AdaTT: Adaptive Task-to-Task Fusion Network for Multi-task Learning in Recommendations. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2023.
- [23] Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. 2017. Focal Loss for Dense Object Detection. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 2017.
- [24] Scott M Lundberg and Su-In Lee. 2017. A Unified Approach to Interpreting Model Predictions. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [25] Jiaqi Ma, Zhe Zhao, Xinyang Yi, Jilin Chen, Lichan Hong, and Ed H Chi. 2018. Modeling Task Relationships in Multi-task Learning with Multi-gate Mixture-of-Experts. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2018.
- [26] Xiao Ma, Liqin Zhao, Guan Huang, Zhi Wang, Zelin Hu, Xiaoqiang Zhu, and Kun Gai. 2018. Entire Space Multi-Task Model: An Effective Approach for Estimating Post-Click Conversion Rate. In *Proceedings of the 41st International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2018.
- [27] Fernando Martinez-Plumed and others. 2023. Improving the Predictive Accuracy of the Cross-selling of Consumer Loans Using Deep Learning Networks. *Annals of Operations Research* (2023).
- [28] Ali Met, Deniz Korkmaz, and Kerem Tuna. 2024. Product Recommendation System With Machine Learning Algorithms for SME Banking. *International Journal of Intelligent Systems* 2024, (2024).
- [29] National Assembly of Korea. 2024. Act on the Development of Artificial Intelligence and Establishment of Trust (AI Basic Act).
- [30] Huy Nguyen, Nhat Ho, and Alessandro Rinaldo. 2024. Sigmoid Gating is More Sample Efficient than Softmax Gating in Mixture of Experts. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [31] Maximilian Nickel and Douwe Kiela. 2017. Poincaré Embeddings for Learning Hierarchical Representations. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [32] Neha Patki, Roy Wedge, and Kalyan Veeramachaneni. 2016. The Synthetic Data Vault. In *IEEE International Conference on Data Science and Advanced Analytics (DSAA)*, 2016.

- [33] Judea Pearl and Dana Mackenzie. 2018. *The Book of Why: The New Science of Cause and Effect*. Basic Books.
- [34] Judea Pearl. 2009. *Causality: Models, Reasoning, and Inference* (2nd ed.). Cambridge University Press.
- [35] Mark Raasveldt and Hannes Mühleisen. 2019. DuckDB: An Embeddable Analytical Database. (2019).
- [36] Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. 2016. Why Should I Trust You?: Explaining the Predictions of Any Classifier. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016.
- [37] Amr Salih and others. 2023. Examining the Limitations of Post-hoc Explainability Methods in Finance. *Frontiers in Artificial Intelligence* (2023).
- [38] Noam Shazeer, Azalia Mirhoseini, Krzysztof Mahdavi, and others. 2017. Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of-Experts Layer. (2017).
- [39] Xiang-Rong Sheng, Liqin Zhao, Guorui Zhou, Xinyao Ding, Binding Dai, Qiang Luo, Siran Yang, Jingshan Lv, Chi Zhang, Hongbo Deng, and Xiaoliang Zhu. 2021. STAR: One-pass Single-Arch Multi-task Architecture for Cross-Domain CTR Prediction. In *Proceedings of the 30th ACM International Conference on Information and Knowledge Management (CIKM)*, 2021.
- [40] Mukund Sundararajan, Ankur Taly, and Qiqi Yan. 2017. Axiomatic Attribution for Deep Networks. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, 2017.
- [41] Hongyan Tang, Junning Liu, Ming Zhao, and Xudong Gong. 2020. Progressive Layered Extraction (PLE): A Novel Multi-Task Learning (MTL) Model for Personalized Recommendations. In *Proceedings of the 14th ACM Conference on Recommender Systems (RecSys)*, 2020.
- [42] Lei Xu, Maria Skoularidou, Alfredo Cuesta-Infante, and Kalyan Veeramachaneni. 2019. Modeling Tabular Data using Conditional GAN. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [43] Hanqing Zeng and others. 2024. Mixture of Weak and Strong Experts on Graphs. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [44] Zijian Zhang and others. 2024. M3oE: Multi-Domain Multi-Task Mixture-of Experts Recommendation Framework. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2024.
- [45] Xun Zheng, Bryon Aragam, Pradeep Ravikumar, and Eric P Xing. 2018. DAGs with NO TEARS: Continuous Optimization for Structure Learning. (2018).